



# Distributed Systems: Remote Procedure Call (RPC), MapReduce

CS 571: *Operating Systems (Spring 2021)*

Lecture 11

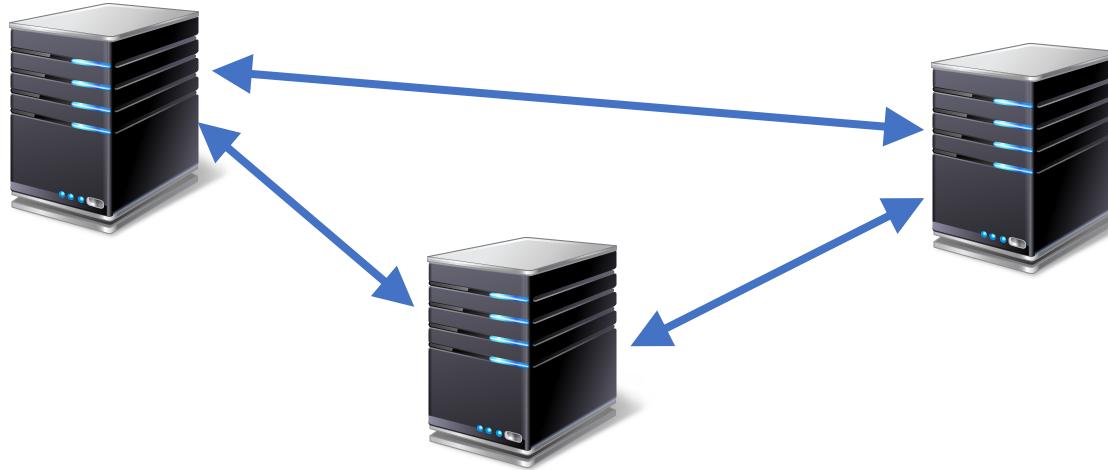
Yue Cheng

Some material taken/derived from:

- Wisconsin CS-537 materials by Remzi Arpacı-Dusseau.

Licensed for use under a Creative Commons Attribution-NonCommercial-ShareAlike 3.0 Unported License.

# What is a distributed system?



- Multiple computers
- Connected by a network
- Doing something together
- A *distributed system* is many cooperating computers that appear to users as a single service

# Today's outline

*How can processes on different cooperating computers exchange information?*

1. Network sockets and raw messages
2. Remote procedure call

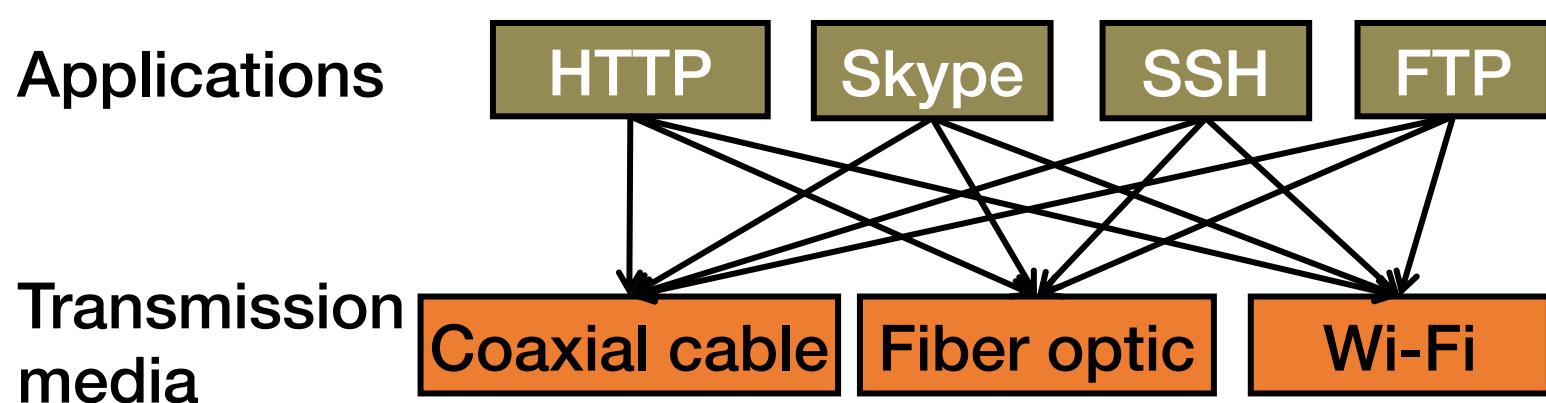
*How can large computing jobs be parallelized?*

1. MapReduce

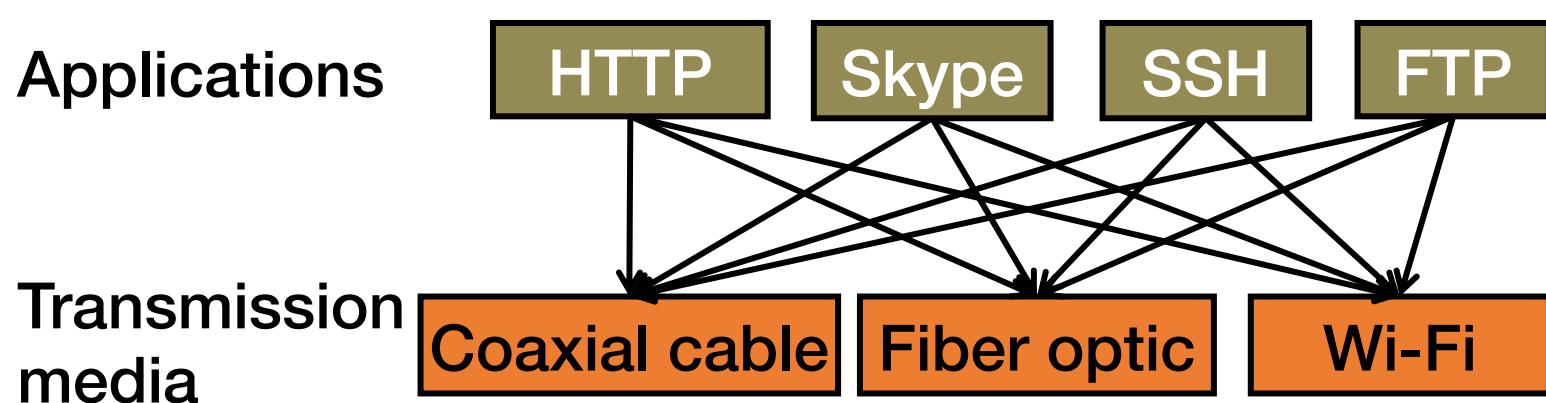
# The problem of communication

- Process on **Host A** wants to talk to process on **Host B**
  - A and B must agree on the meaning of the bits being sent and received at many different levels, including:
    - How many volts is a 0 bit, a 1 bits?
    - How does receiver know which is the last bit?
    - How many bits long is a number?

# The problem of communication

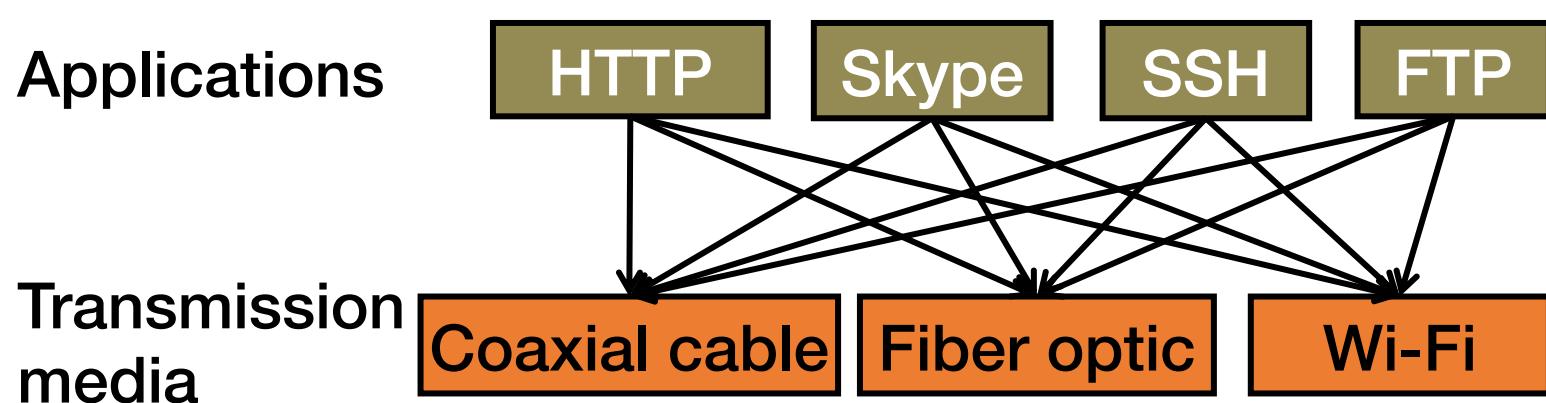


# The problem of communication



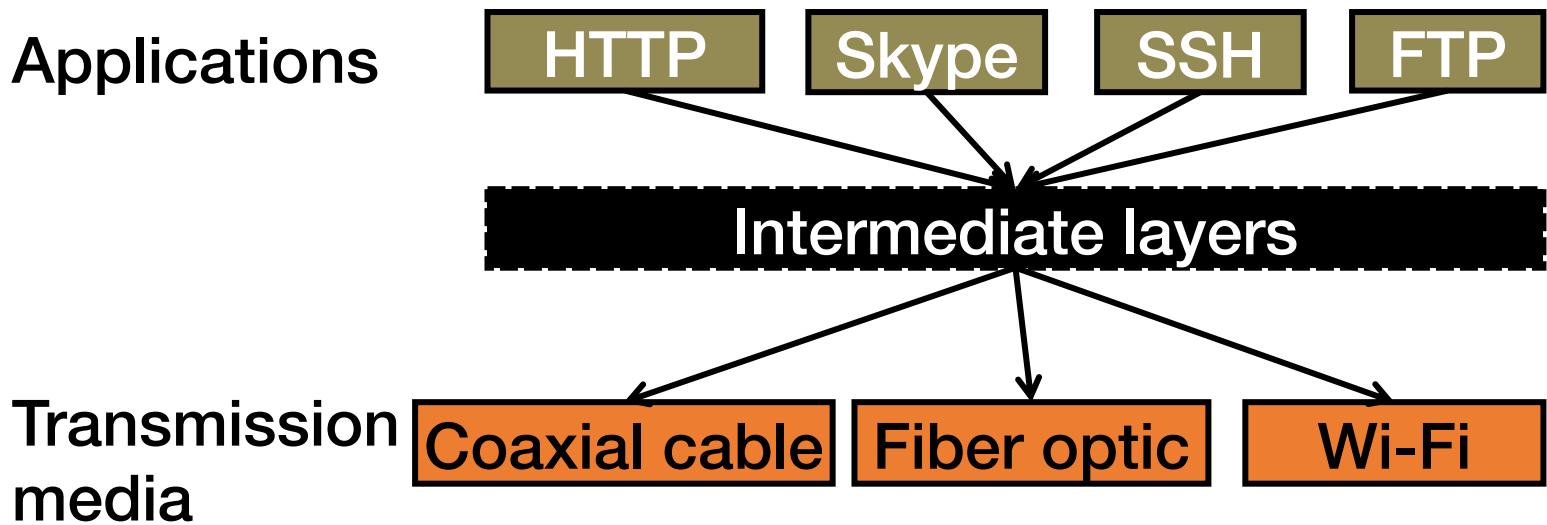
- Re-implement every application for every new underlying transmission medium?
- Change every application on any change to an underlying transmission medium?

# The problem of communication



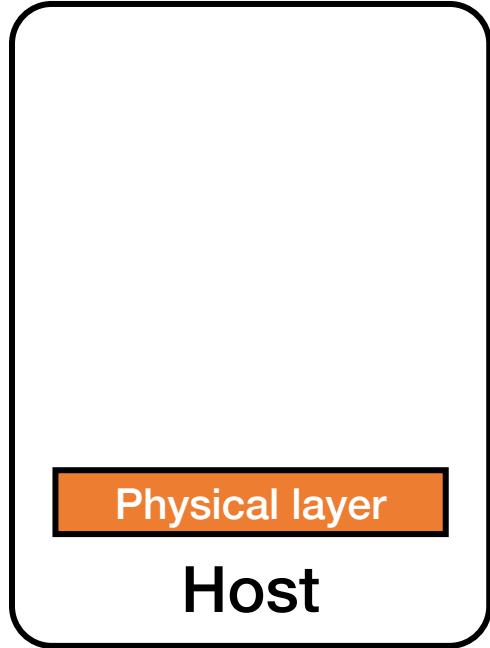
- Re-implement every application for every new underlying transmission medium?
- Change every application on any change to an underlying transmission medium?
- No! But how does the Internet design avoid this?

# Solution: Layering



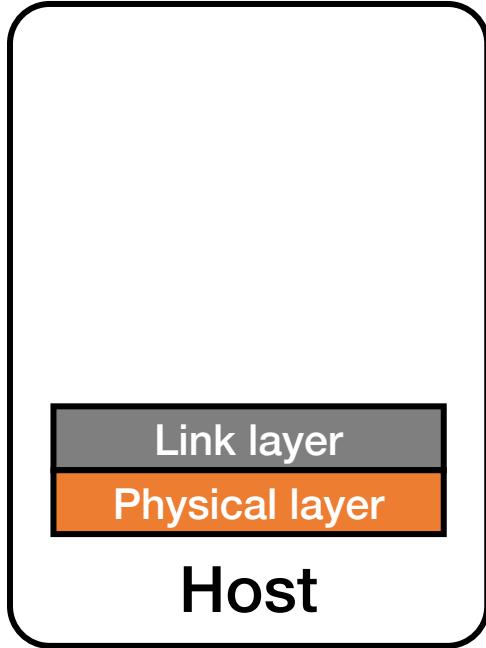
- Intermediate **layers** provide a set of abstractions for applications and media
- New applications or media need only implement for intermediate layer's interface

# Layering in the Internet



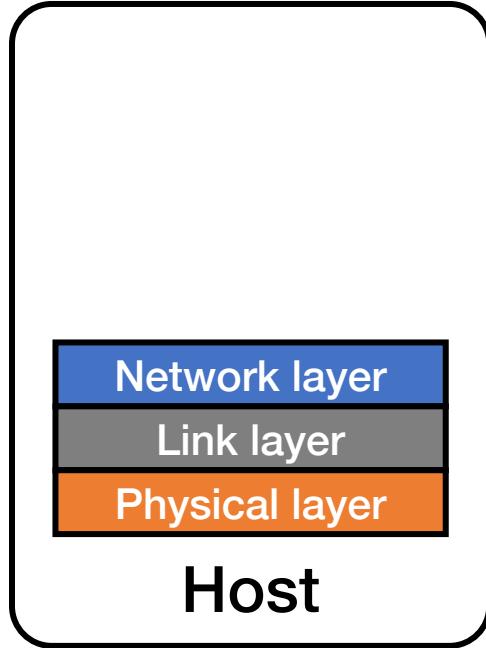
- **Physical:** Moves bits between two hosts connected by a physical link

# Layering in the Internet



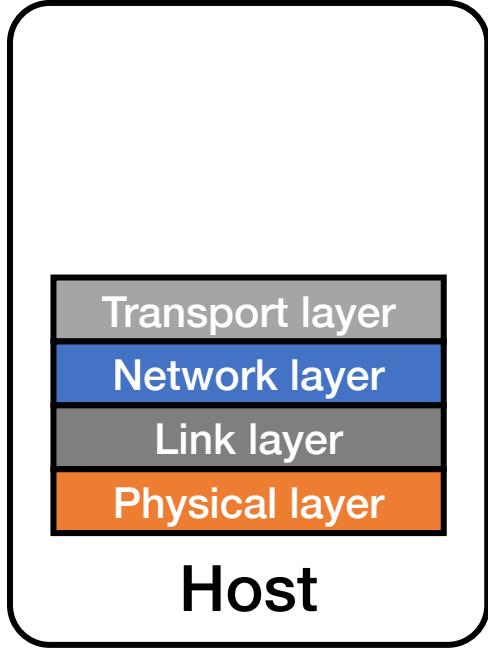
- **Link:** Enables end hosts to exchange atomic messages with each other
- **Physical:** Moves bits between two hosts connected by a physical link

# Layering in the Internet



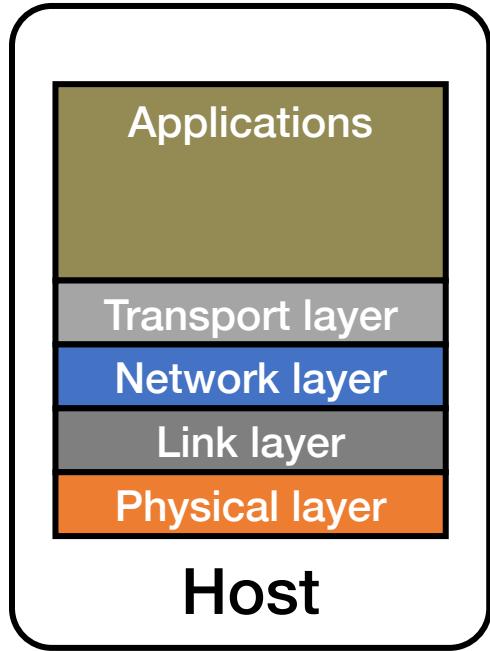
- **Network:** Deliver packets to destinations on other (heterogeneous) networks
- **Link:** Enables end hosts to exchange atomic messages with each other
- **Physical:** Moves bits between two hosts connected by a physical link

# Layering in the Internet



- **Transport:** Provide end-to-end communication between processes on different hosts
- **Network:** Deliver packets to destinations on other (heterogeneous) networks
- **Link:** Enables end hosts to exchange atomic messages with each other
- **Physical:** Moves bits between two hosts connected by a physical link

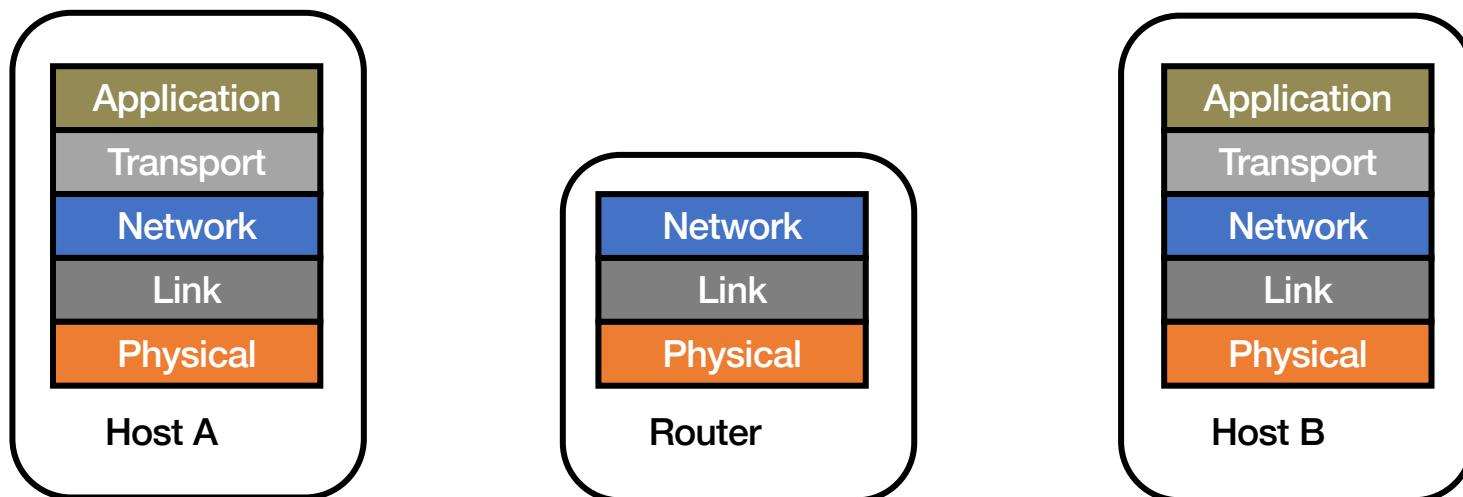
# Layering in the Internet



- **Transport:** Provide end-to-end communication between processes on different hosts
- **Network:** Deliver packets to destinations on other (heterogeneous) networks
- **Link:** Enables end hosts to exchange atomic messages with each other
- **Physical:** Moves bits between two hosts connected by a physical link

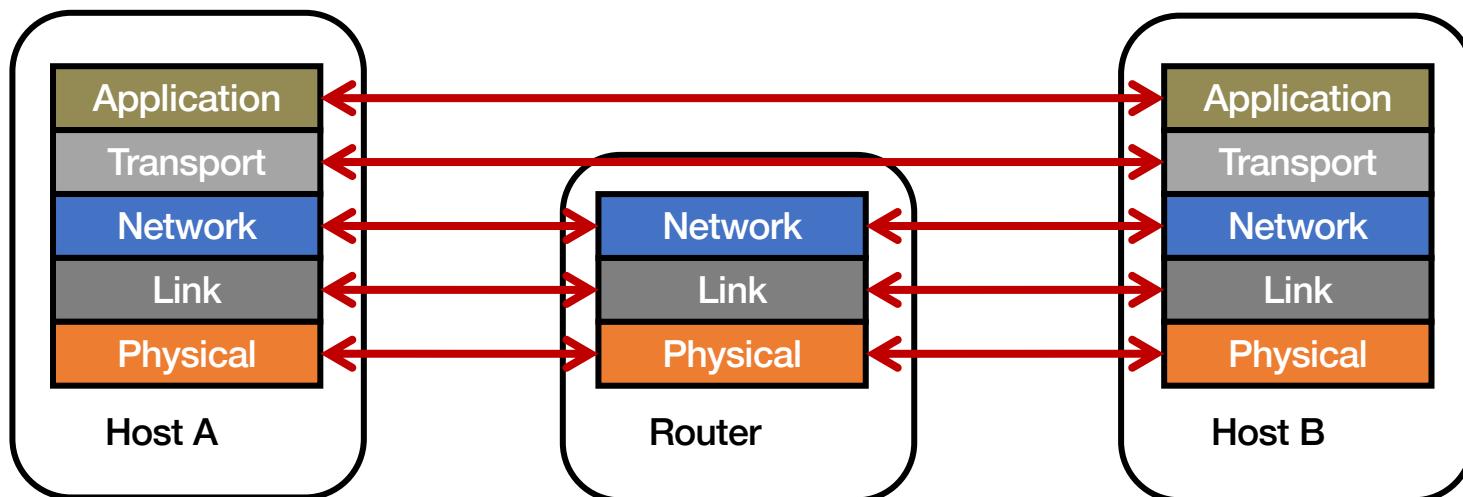
# Logical communication between layers

- How to forge agreement on the meaning of the bits exchanged between two hosts?



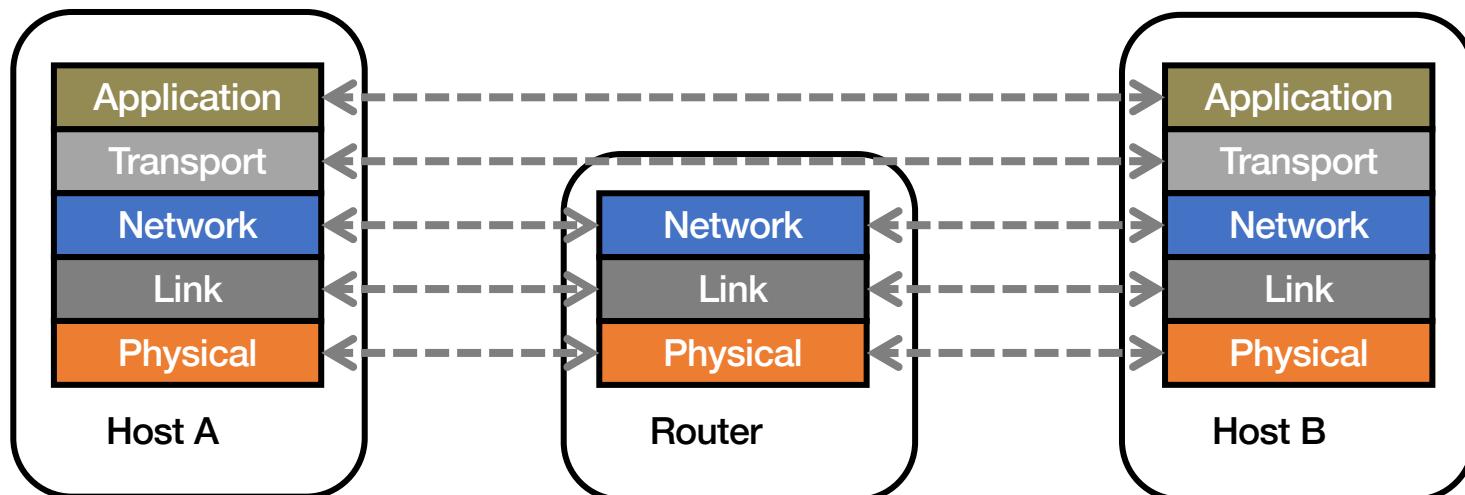
# Logical communication between layers

- How to forge agreement on the meaning of the bits exchanged between two hosts?
- **Protocol:** Rules that govern the format, contents, and meaning of messages
  - Each layer on a host interacts with its peer host's corresponding layer via the **protocol interface**



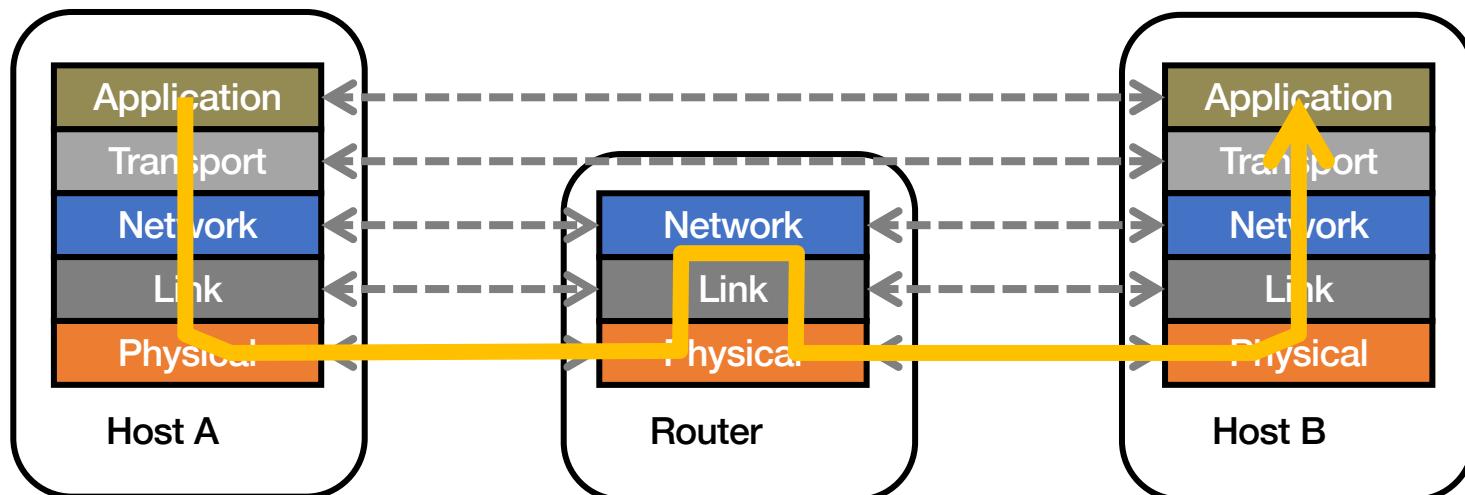
# Physical communication

- Communication goes down to the **physical network**
- Then from **network** peer to peer
- Then up to the relevant application



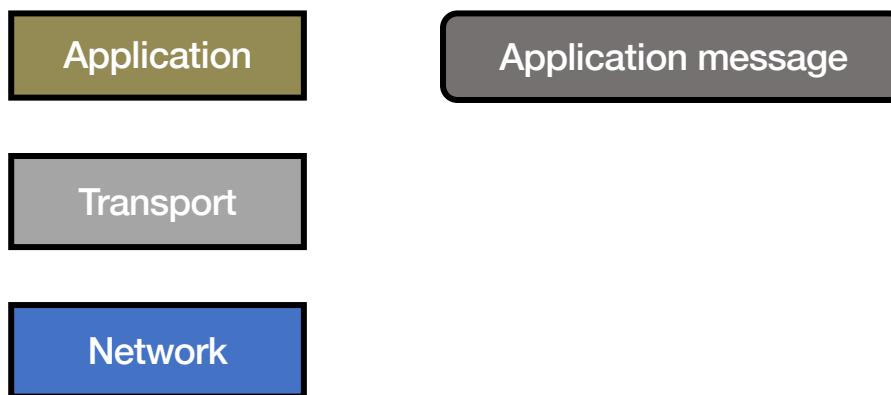
# Physical communication

- Communication goes down to the **physical network**
- Then from **network** peer to peer
- Then up to the relevant application



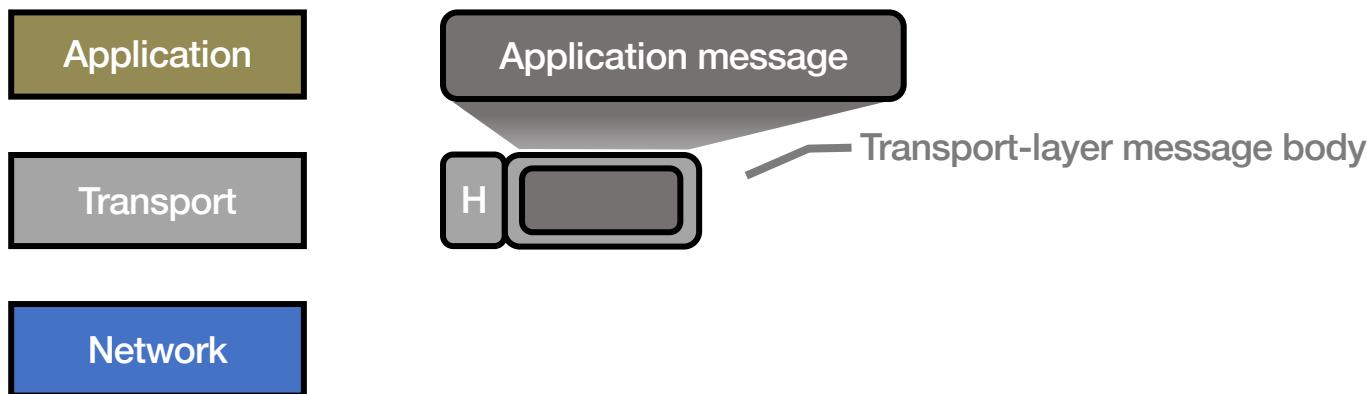
# Communication between layers

- How do peer protocols coordinate with each other?
- Layer attaches its own header (H) to communicate with peer
  - Higher layers' headers, data **encapsulated** inside message
    - Lower layers don't generally inspect higher layers' headers



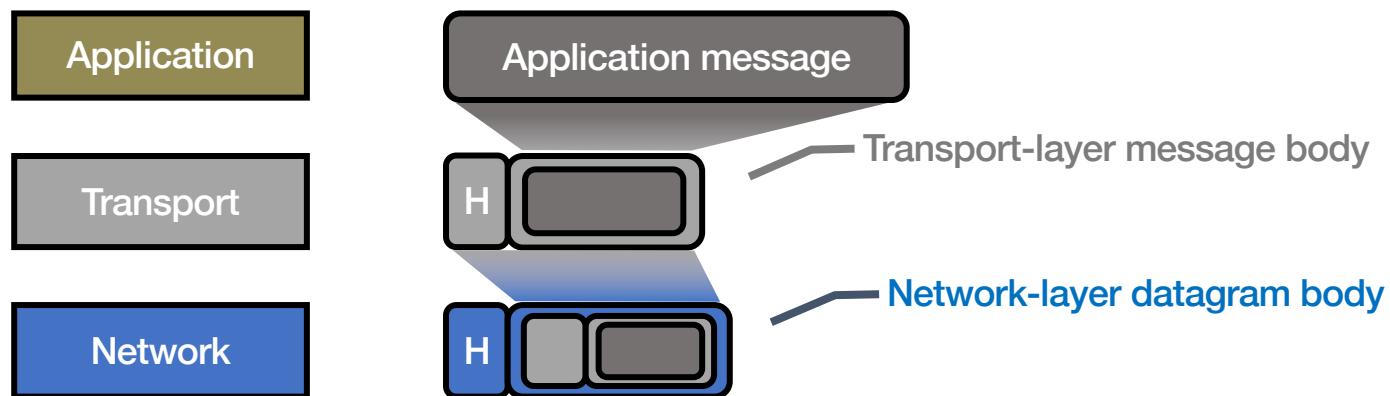
# Communication between layers

- How do peer protocols coordinate with each other?
- Layer attaches its own header (H) to communicate with peer
  - Higher layers' headers, data **encapsulated** inside message
    - Lower layers don't generally inspect higher layers' headers



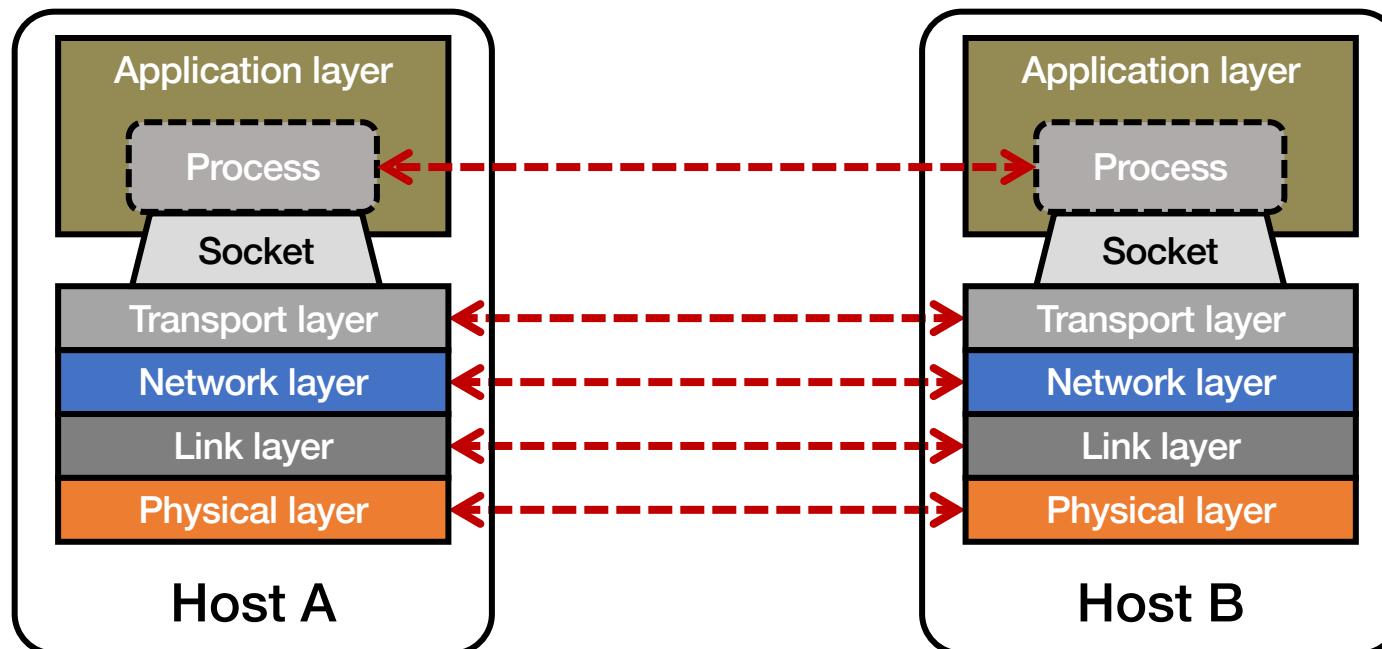
# Communication between layers

- How do peer protocols coordinate with each other?
- Layer attaches its own header (H) to communicate with peer
  - Higher layers' headers, data **encapsulated** inside message
    - Lower layers don't generally inspect higher layers' headers



# Network socket-based communication

- **Socket:** The interface the OS provides to the network
  - Provides inter-process explicit message exchange
- Can build distributed systems atop sockets: `send()`, `recv()`
  - e.g.: **put(key, value)** → message



# Network sockets: Summary

- Principle of transparency: Hide that resource is physically distributed across multiple computers
  - Access resource same way as locally
  - Users can't tell where resource is physically located

Network sockets provide apps with point-to-point communication between processes

- **put(key,value)** → message with sockets?

```
// Create a socket for the client
if ((sockfd = socket (AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("Socket creation");
    exit(2);
}

// Set server address and port
memset(&servaddr, 0, sizeof(servaddr));
servaddr.sin_family = AF_INET;
servaddr.sin_addr.s_addr = inet_addr(argv[1]);
servaddr.sin_port = htons(SERV_PORT); // to big-endian

// Establish TCP connection
if (connect(sockfd, (struct sockaddr *) &servaddr,
            sizeof(servaddr)) < 0) {
    perror("Connect to server");
    exit(3);
}

// Transmit the data over the TCP connection
send(sockfd, buf, strlen(buf), 0);
```

```
// Create a socket for the client
if ((sockfd = socket (AF_INET, SOCK_STREAM, 0)) < 0) {
    perror("Socket creation");
    exit(2);
}

// Set server address and port
memset(&servaddr, 0, sizeof(servaddr));
servaddr.sin_family = AF_INET;
servaddr.sin_addr.s_addr = inet_addr(argv[1]);
servaddr.sin_port = htons(SERV_PORT); // to big-endian

// Establish TCP connection
if (connect(sockfd, (struct sockaddr *) &servaddr,
            sizeof(servaddr)) < 0) {
    perror("Connect to server");
    exit(3);
}

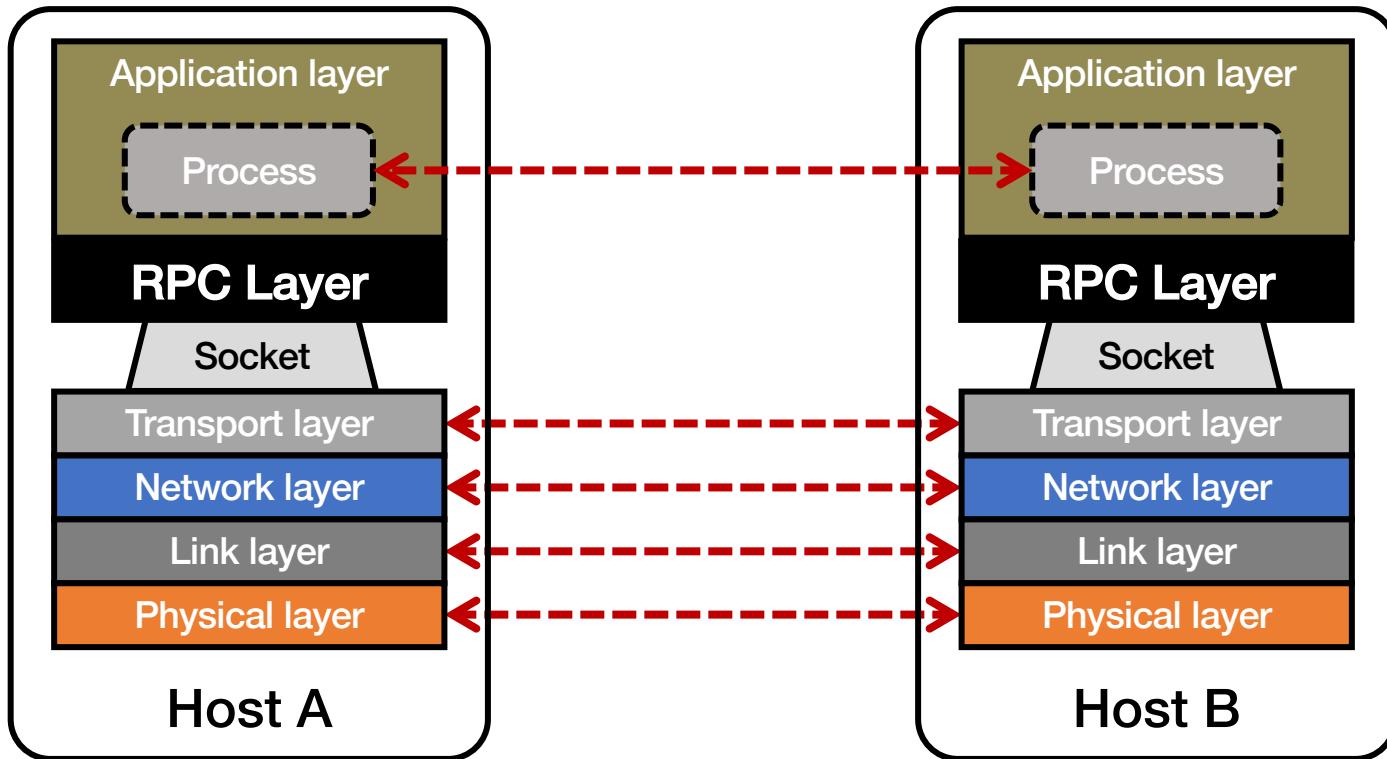
// Transmit the data over the TCP connection
send(sockfd, buf, strlen(buf), 0);
```

Sockets don't provide transparency

# Takeaway: Socket programming still not ideal (great)

- Lots for the programmer to deal with every time
  - How to separate different requests on the same connection?
  - How to write bytes to the network / read bytes from the network?
    - What if Host A's process is written in Go and Host B's process is in C++?
  - What to do with those bytes?
- Still pretty **painful**... Have to worry a lot about the network

# Solution: Another layer!



# Today's outline

*How can processes on different cooperating computers exchange information?*

1. Network sockets and raw messages
2. Remote procedure call

*How can large computing jobs be parallelized?*

3. MapReduce

# Motivation: Why RPC?

- The typical programmer is trained to write single-threaded code that runs in one place
- Goal: Easy-to-program network communication that makes client-server communication **transparent**
  - Retains the “feel” of writing centralized code
    - Programmer needn’t think about the network
- Programming Labs use Go RPC

# What's the goal of RPC?

- Within a single program, running in a single process, recall the well-known notion of a **procedure call**:
  - **Caller** pushes arguments onto stack,
    - jumps to address of **callee** function
  - **Callee** reads arguments from stack,
    - executes, puts return value in register,
    - returns to next instruction in caller

# What's the goal of RPC?

- Within a single program, running in a single process, recall the well-known notion of a **procedure call**:
  - **Caller** pushes arguments onto stack,
    - jumps to address of **callee** function
  - **Callee** reads arguments from stack,
    - executes, puts return value in register,
    - returns to next instruction in caller

**RPC's Goal:** make communication appear like a local procedure call: transparency for procedure calls – way less painful than sockets...

# RPC issues

## 1. Heterogeneity

- Client needs to rendezvous with the server
- Server must dispatch to the required function
  - What if server is different type of machine?

# RPC issues

## 1. Heterogeneity

- Client needs to rendezvous with the server
- Server must dispatch to the required function
  - What if server is different type of machine?

## 2. Failure

- What if messages get **dropped**?
- What if client, server, or network **fails**?

# RPC issues

## 1. Heterogeneity

- Client needs to rendezvous with the server
- Server must dispatch to the required function
  - What if server is different type of machine?

## 2. Failure

- What if messages get **dropped**?
- What if client, server, or network **fails**?

## 3. Performance

- Procedure call takes takes  $\approx 10$  cycles  $\approx 3$  ns
- RPC in a data center takes  $\approx 10 \mu\text{s}$  ( $10^3 \times$  slower)
  - In the wide area, typically  $10^6 \times$  slower

# Problem: Differences in data representation

- Not an issue for local procedure calls
- For a remote procedure call, a remote machine may:
  - Run process written in a **different language**
  - Represent data types using **different sizes**
  - Use a **different byte ordering** (endianness)
  - Represent floating point numbers **differently**
  - Have **different data alignment** requirements
    - e.g., 4-byte type begins only on 4-byte memory boundary

# Problem: Differences in programming support

- Language support **varies**:
  - Many programming languages have **no inbuilt** way of extracting values from complex types
    - C, C++
    - Effectively need sockets glue code underneath
  - Some languages have support that enables RPC
    - Python, Go
    - Exploit type system for some help

# Solution: Interface Description Language

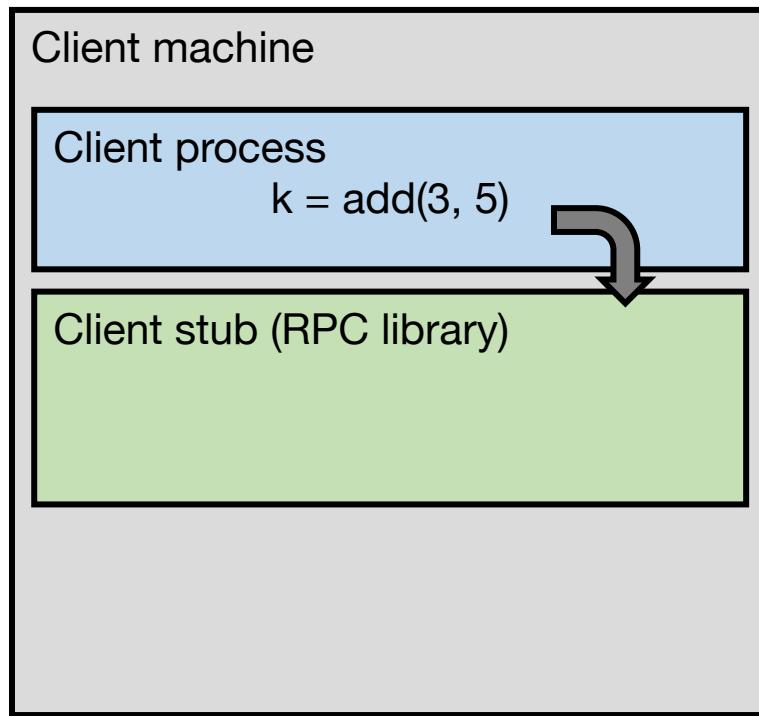
- Mechanism to pass procedure parameters and return values in a machine-independent way
- Programmer may write an **interface description** in the IDL
  - Defines API for procedure calls: names, parameter/return types

# Solution: Interface Description Language

- Mechanism to pass procedure parameters and return values in a machine-independent way
- Programmer may write an **interface description** in the IDL
  - Defines API for procedure calls: names, parameter/return types
- Then runs an **IDL compiler** which generates:
  - Code to **marshal** (convert) native data types into machine-independent byte streams
    - And vice-versa, called **unmarshaling**
  - Client stub: Forwards local procedure call as a request to server
  - Server stub: Dispatches RPC to its implementation

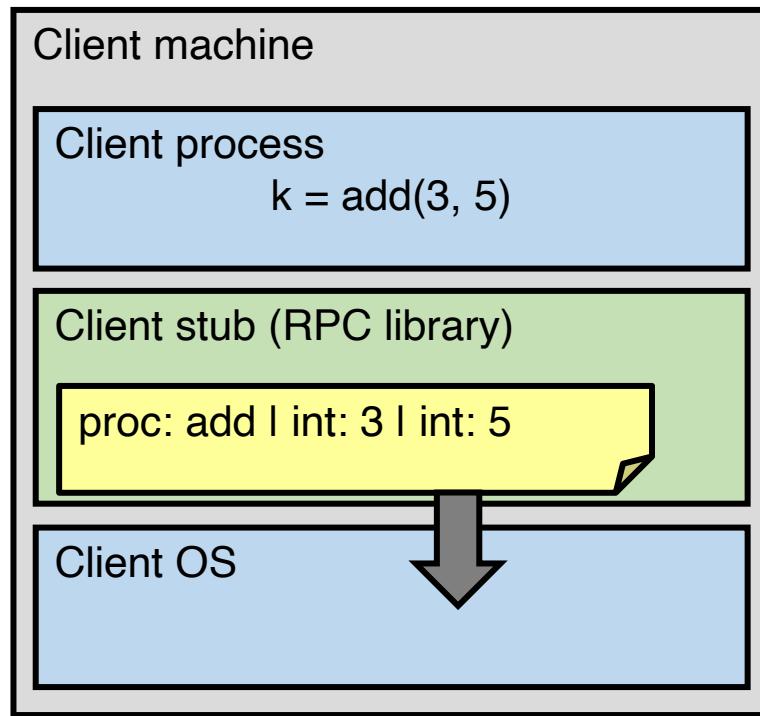
# A day in the life of an RPC

1. Client calls stub function (pushes parameters onto stack)



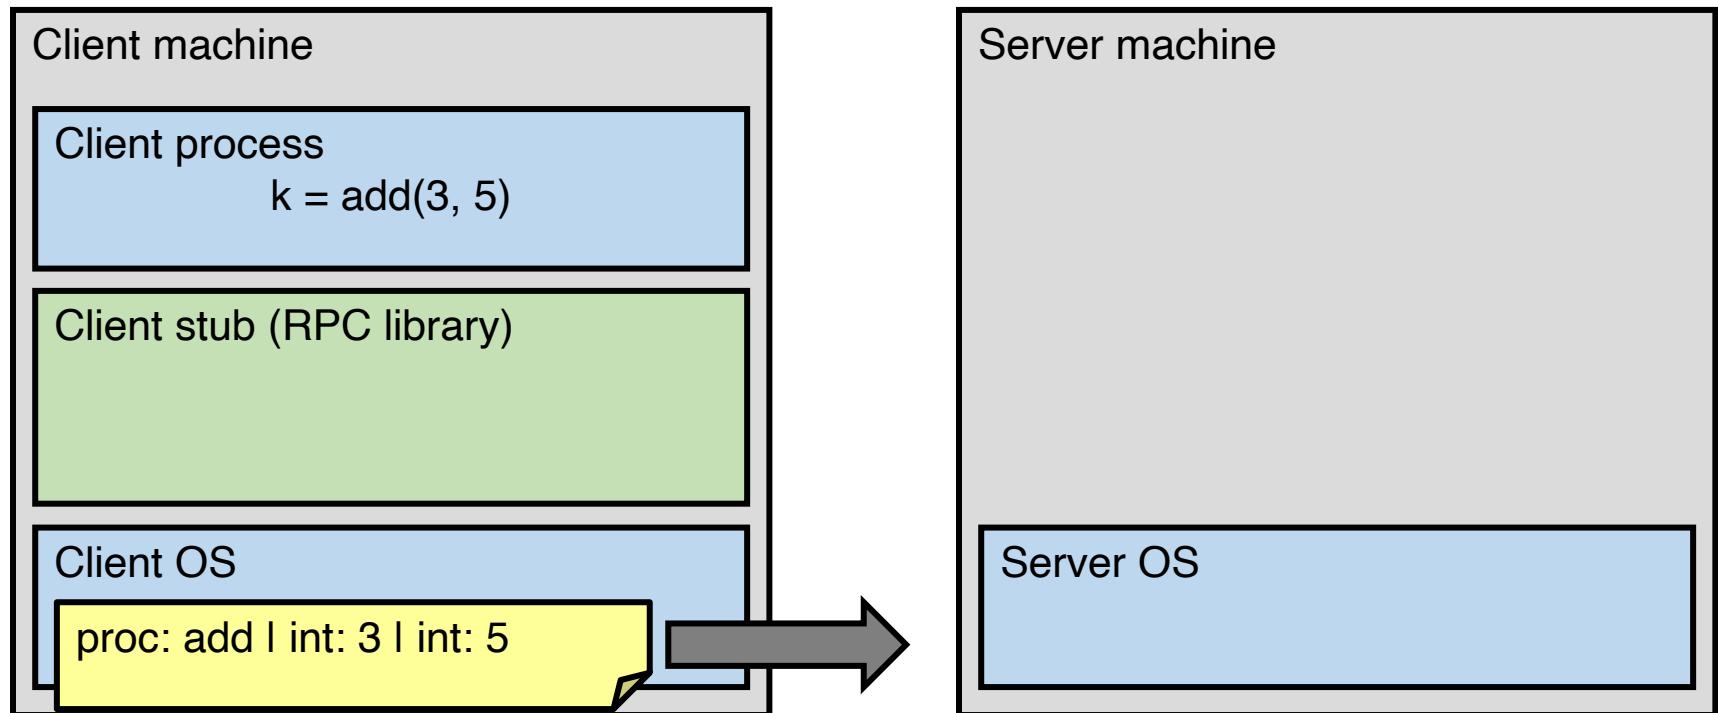
# A day in the life of an RPC

1. Client calls stub function (pushes parameters onto stack)
2. Stub marshals parameters to a network message



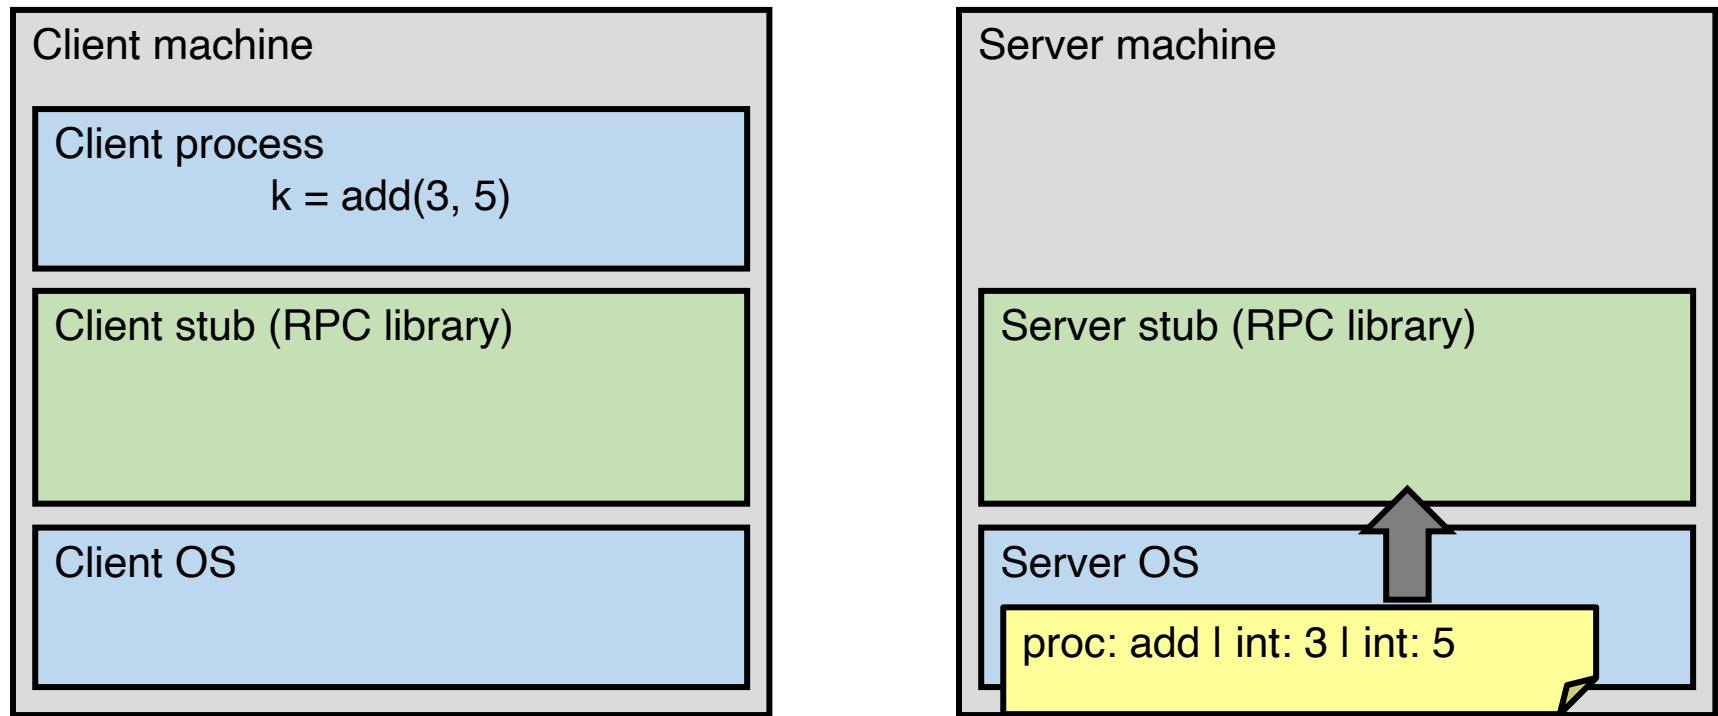
# A day in the life of an RPC

2. Stub marshals parameters to a network message
3. OS sends a network message to the server



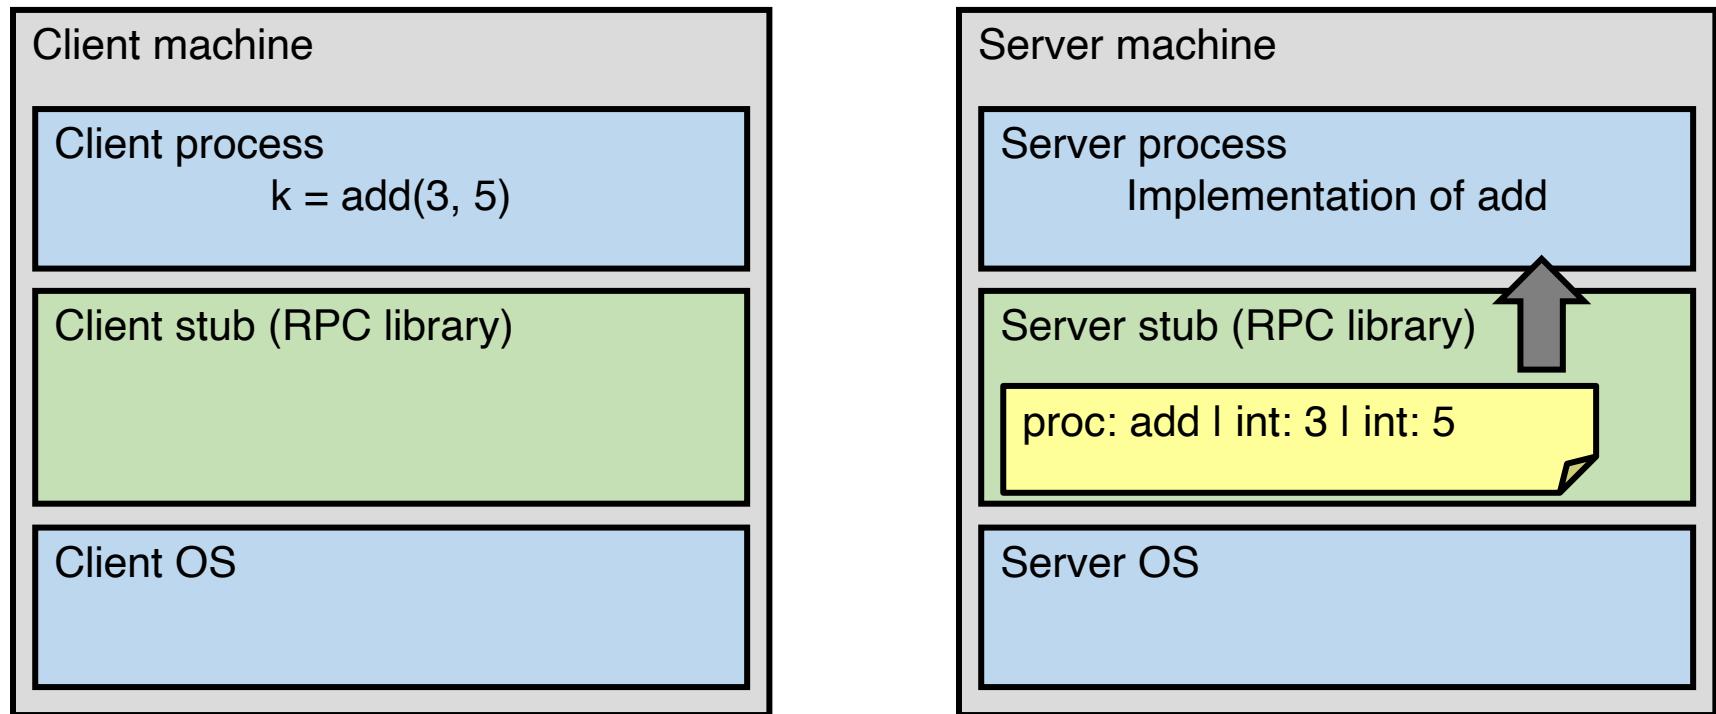
# A day in the life of an RPC

3. OS sends a network message to the server
4. Server OS receives message, sends it up to stub



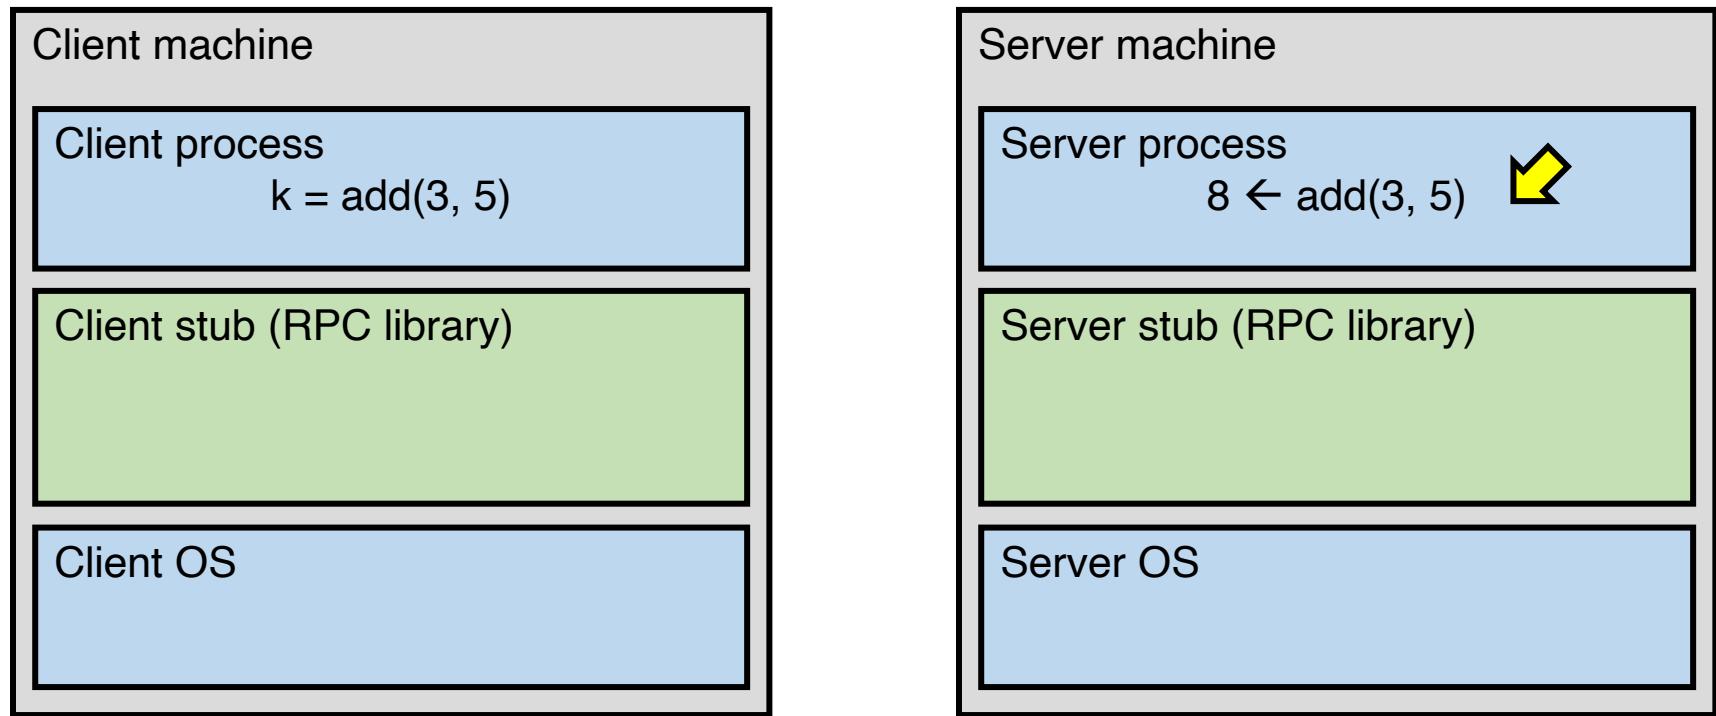
# A day in the life of an RPC

4. Server OS receives message, sends it up to stub
5. Server stub unmarshals params, calls server function



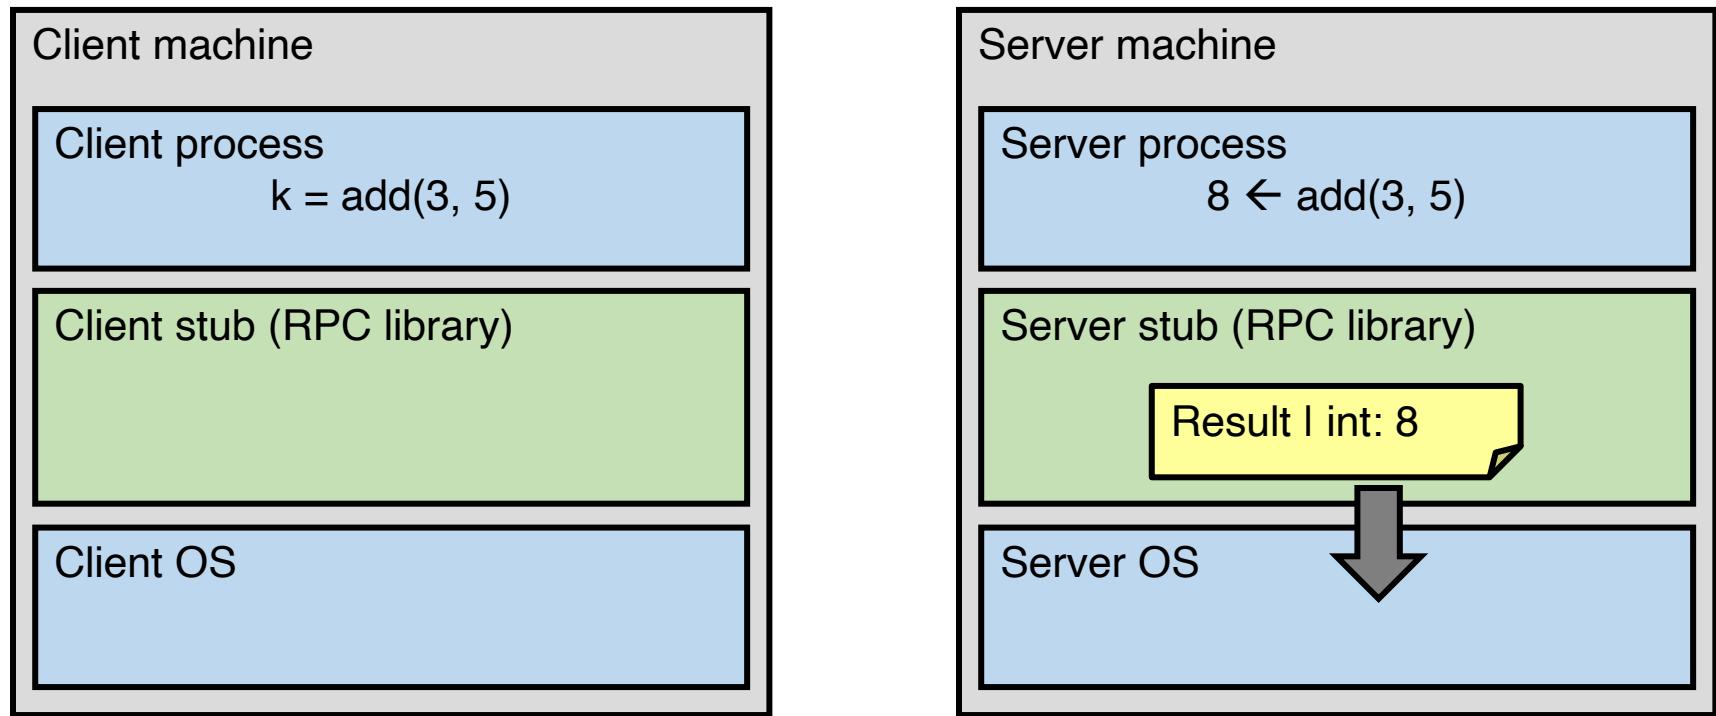
# A day in the life of an RPC

5. Server stub unmarshals params, calls server function
6. Server function runs, returns a value



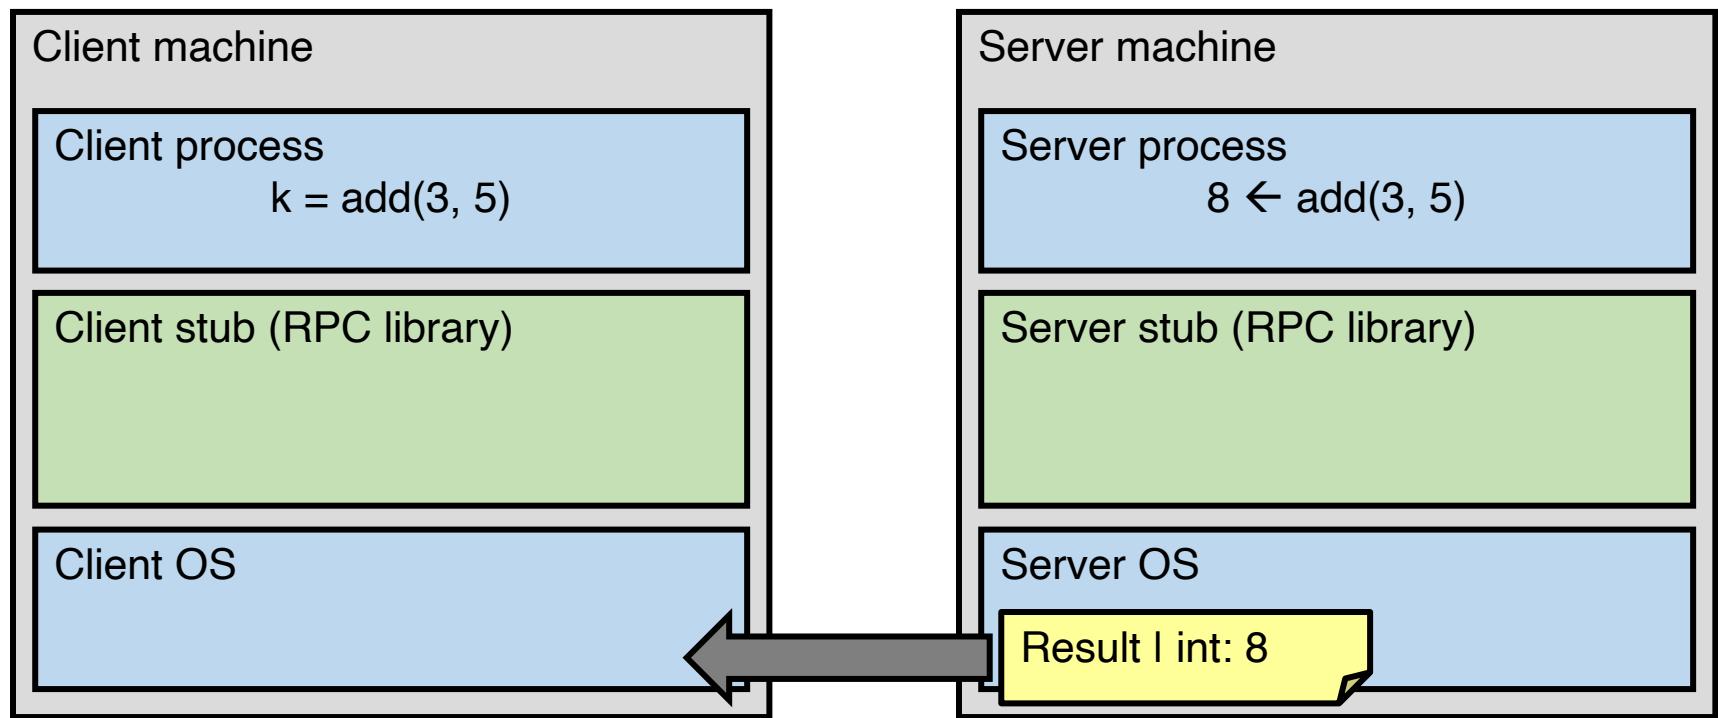
# A day in the life of an RPC

6. Server function runs, returns a value
7. Server stub marshals the return value, sends message



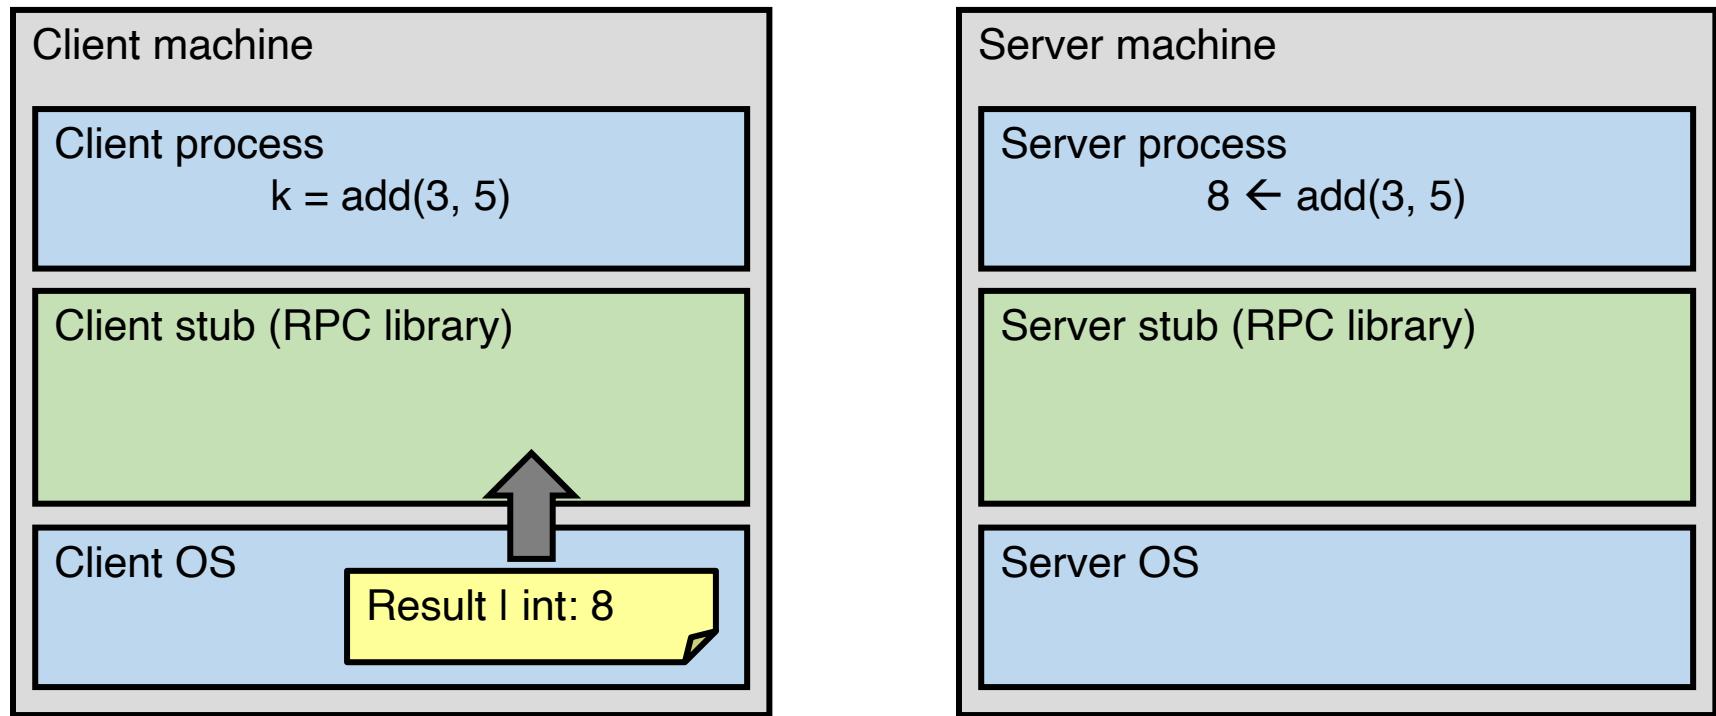
# A day in the life of an RPC

7. Server stub marshals the return value, sends message
8. Server OS sends the reply back across the network



# A day in the life of an RPC

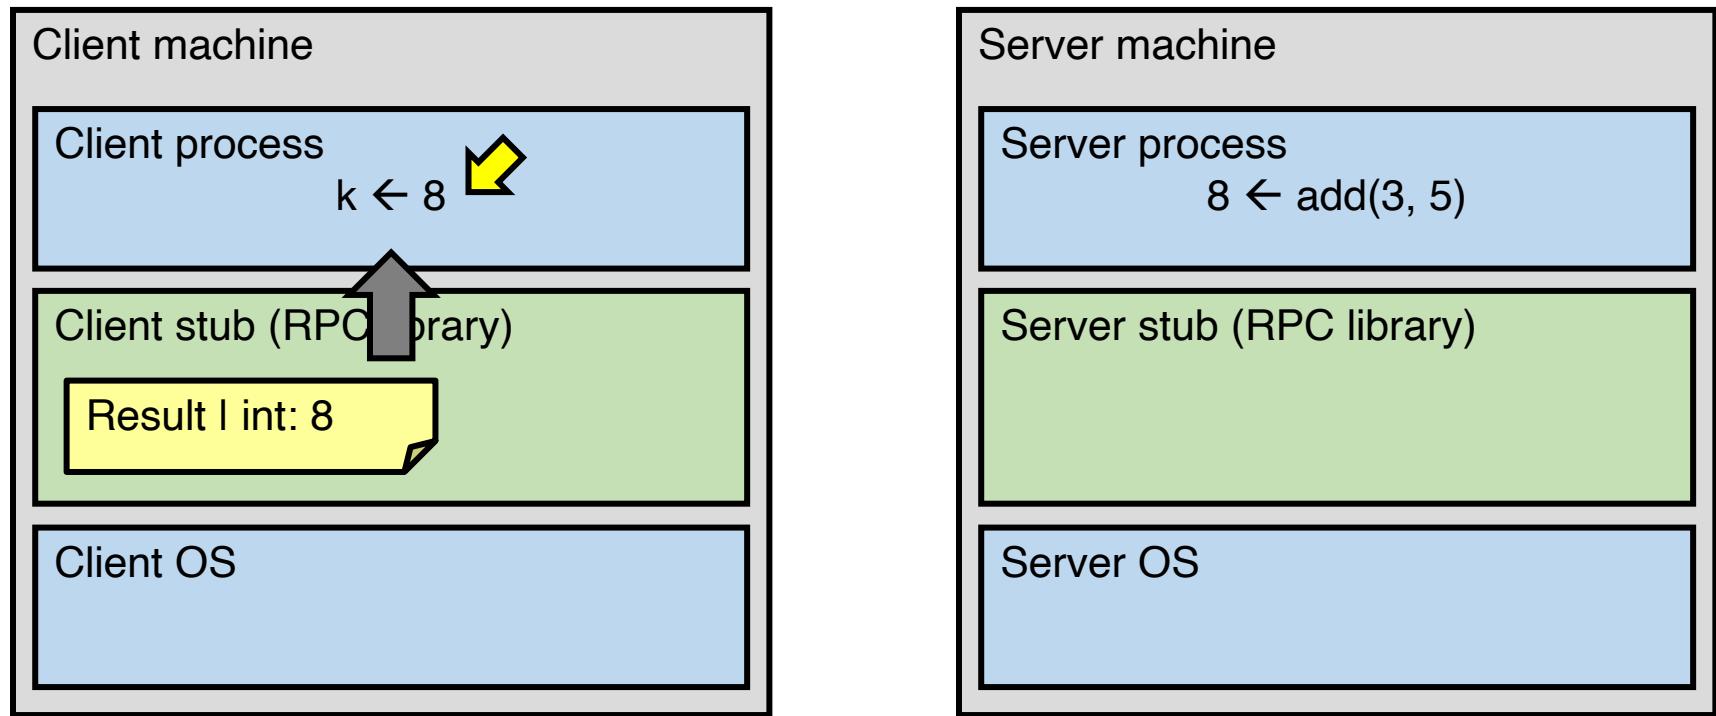
8. Server OS sends the reply back across the network
9. Client OS receives the reply and passes up to stub



# A day in the life of an RPC

9. Client OS receives the reply and passes up to stub

10. Client stub unmarshals return value, returns to client



# The server stub is really two parts

- Dispatcher
  - Receives a client's RPC request
    - Identifies appropriate server-side method to invoke
- Skeleton
  - Unmarshals parameters to server-native types
  - Calls the local server procedure
  - Marshals the response, sends it back to the dispatcher
- All this is hidden from the programmer
  - Dispatcher and skeleton may be integrated
    - Depends on implementation

# Today's outline

*How can processes on different cooperating computers exchange information?*

1. Network sockets and raw messages
2. Remote procedure call

*How can large computing jobs be parallelized?*

3. MapReduce

## Applications

Web  
apps

Data  
processing

Data  
storage

Emerging  
apps?

## Resource management

Compute  
resources

Memory  
resources

Storage  
resources

Network  
resources



## Datacenter infrastructure



## Applications

Web  
apps

Data  
processing

Data  
storage

Emerging  
apps?

## Resource management

Compute

Memory

Storage

Network

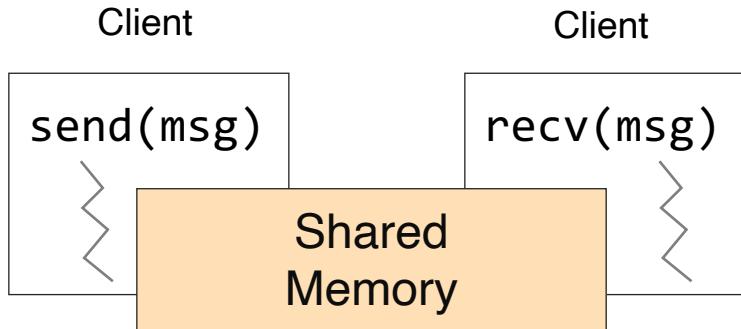
**Question:** How to program these many computers?



Datacenter architecture

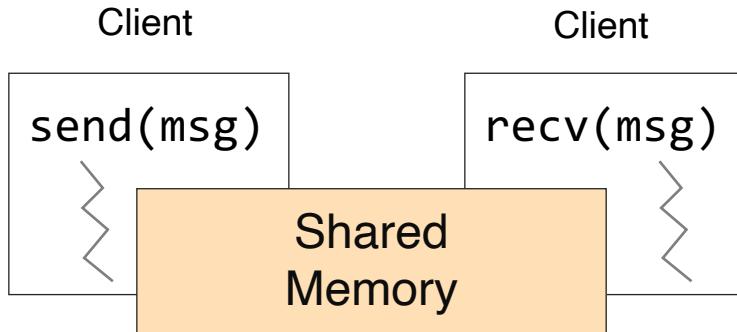


# Shared memory

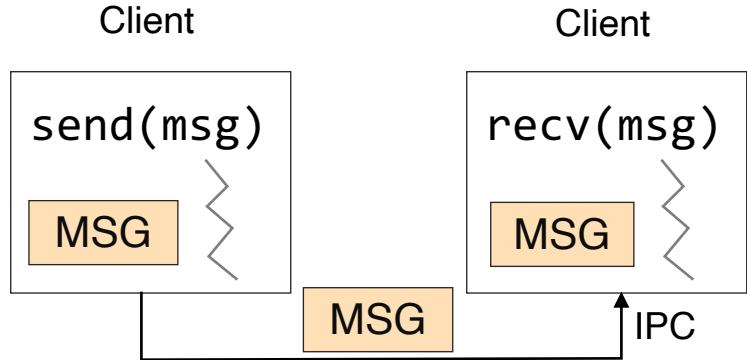


- Shared memory: allows multiple processes to share data via memory
- Applications must locate and map shared memory regions to exchange data

# Shared memory vs. Message passing



- Shared memory: all multiple processes share data via memory
- Applications must locate and map shared memory regions to exchange data



- Message passing: exchange data explicitly via IPC
- Application developers define protocol and exchanging format, number of participants, and each exchange

# Shared memory vs. Message passing

- Easy to program; just like a single multi-threaded machines
- Hard to write high perf. apps:
  - Cannot control which data is local or remote (remote mem. access much slower)
- Hard to mask failures
- Message passing: can write very high perf. apps
- Hard to write apps:
  - Need to manually decompose the app, and move data
  - Need to manually handle failures

# Shared memory: Pthread

- A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization
- API specifies behavior of the thread library, implementation is up to development of the library
- Common in UNIX (e.g., Linux) OSes

# Shared memory: Pthread

```
void *myThreadFun(void *vargp) {
    sleep(1);
    printf("Hello world\n");
    return NULL;
}

int main() {
    pthread_t thread_id_1, thread_id_2;
    pthread_create(&thread_id_1, NULL, myThreadFun, NULL);
    pthread_create(&thread_id_2, NULL, myThreadFun, NULL);
    pthread_join(thread_id_1, NULL);
    pthread_join(thread_id_2, NULL);
    exit(0);
}
```

# Message passing: MPI

- MPI – Message Passing Interface
  - Library standard defined by a committee of vendors, implementers, and parallel programmers
  - Used to create parallel programs based on message passing
- Portable: one standard, many implementations
  - Available on almost all parallel machines in C and Fortran
  - De facto standard platform for the HPC community

# Message passing: MPI

```
int main(int argc, char **argv) {
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, *world_rank);

    // Print off a hello world message
    printf("Hello world from rank %d out of %d processors\n",
           world_rank, world_size);

    // Finalize the MPI environment
    MPI_Finalize();
}
```

# Message passing: MPI

```
mpirun -n 4 -f host_file ./mpi_hello_world
```

```
int main(int argc, char **argv) {
    MPI_Init(NULL, NULL);

    // Get the number of processes
    int world_size;
    MPI_Comm_size(MPI_COMM_WORLD, &world_size);

    // Get the rank of the process
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, *world_rank);

    // Print off a hello world message
    printf("Hello world from rank %d out of %d processors\n",
           world_rank, world_size);

    // Finalize the MPI environment
    MPI_Finalize();
}
```

# MapReduce

# The big picture (motivation)

- Datasets are **too big** to process using a single computer

# The big picture (motivation)

- Datasets are **too big** to process using a single computer
- Good parallel processing engines are **rare** (back then in the late 90s)

# The big picture (motivation)

- Datasets are **too big** to process using a single computer
- Good parallel processing engines are **rare** (back then in the late 90s)
- Want a parallel processing framework that:
  - is **general** (works for many problems)
  - is **easy to use** (no locks, no need to explicitly handle communication, no race conditions)
  - can **automatically parallelize** tasks
  - can **automatically handle** machine failures

# Context (Google circa 2000)

- Starting to deal with **massive** datasets
- But also addicted to cheap, unreliable hardware
  - Young company, expensive hardware not practical
- Only a few expert programmers can write distributed programs to process them
  - Scale so large jobs can complete before failures



# Context (Google circa 2000)

- Starting to deal with **massive** datasets
- But also addicted to cheap, unreliable hardware
  - Young company, expensive hardware not practical
- Only a few expert programmers can write distributed programs to process them
  - Scale so large jobs can complete before failures
- **Key question:** how can every Google engineer be imbued with the ability to write **parallel**, **scalable**, **distributed**, **fault-tolerant** code?
- **Solution:** **abstract out** the redundant parts
- **Restriction:** relies on job semantics, so restricts which problems it works for

# Application: Word Count

```
cat data.txt  
| tr -s '[:punct:][:space:]' '\n'  
| sort | uniq -c
```

```
SELECT count(word), word FROM data  
GROUP BY word
```

# Deal with multiple files?

1. Compute word counts from individual files

# Deal with multiple files?

1. Compute word counts from individual files
2. Then merge intermediate output

# Deal with multiple files?

1. Compute word counts from individual files
2. Then merge intermediate output
3. Compute word count on merged outputs

# What if the data is too big to fit in one computer?

1. In parallel, send to worker:
  - Compute word counts from individual files
  - Collect results, wait until all finished

# What if the data is too big to fit in one computer?

1. In parallel, send to worker:
  - Compute word counts from individual files
  - Collect results, wait until all finished
2. Then merge intermediate output

# What if the data is too big to fit in one computer?

1. In parallel, send to worker:
  - Compute word counts from individual files
  - Collect results, wait until all finished
2. Then merge intermediate output
3. Compute word count on merged intermediates

# MapReduce: Programming interface

- $\text{map}(\text{k1}, \text{v1}) \rightarrow \text{list}(\text{k2}, \text{v2})$ 
  - Apply function to  $(\text{k1}, \text{v1})$  pair and produce set of intermediate pairs  $(\text{k2}, \text{v2})$
- $\text{reduce}(\text{k2}, \text{list}(\text{v2})) \rightarrow \text{list}(\text{k3}, \text{v3})$ 
  - Apply aggregation (reduce) function to values
  - Output results

# MapReduce: Word Count

```
map(key, value):
```

```
    for each word w in value:
```

```
        EmitIntermediate(w, "1");
```

```
reduce(key, values):
```

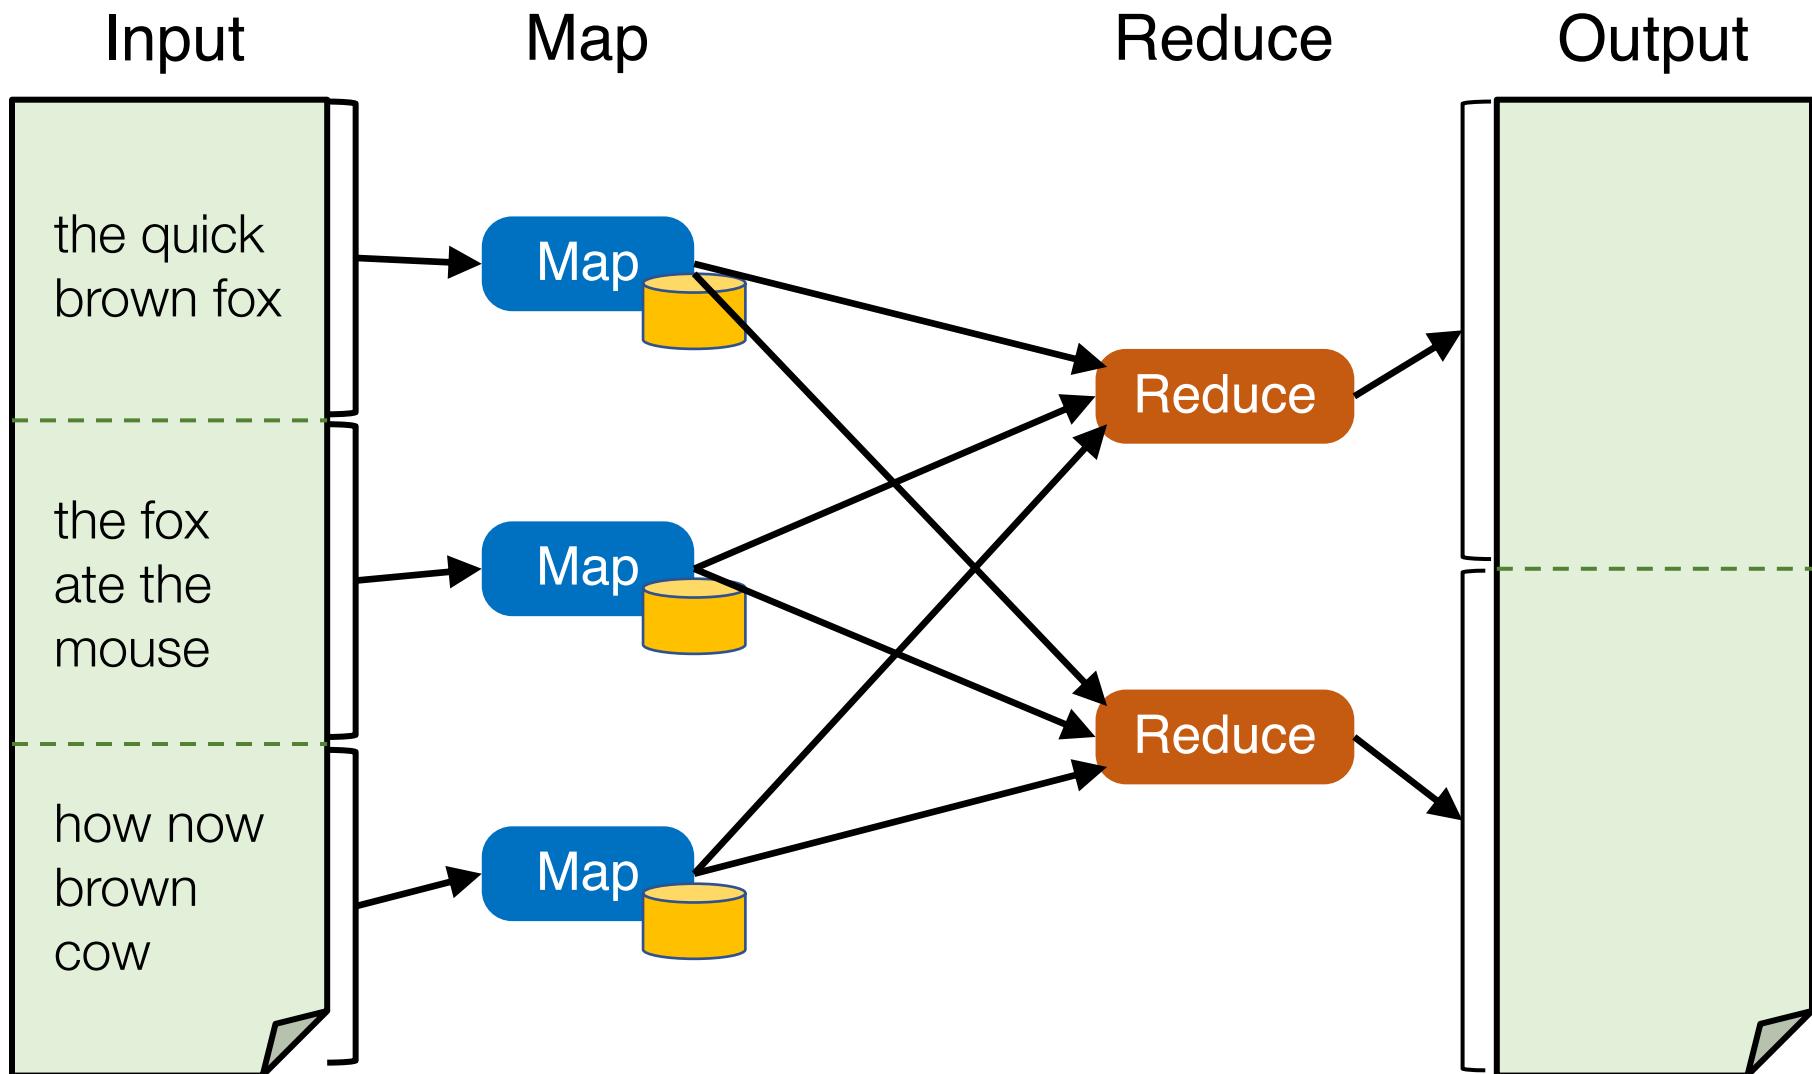
```
    int result = 0;
```

```
    for each v in values:
```

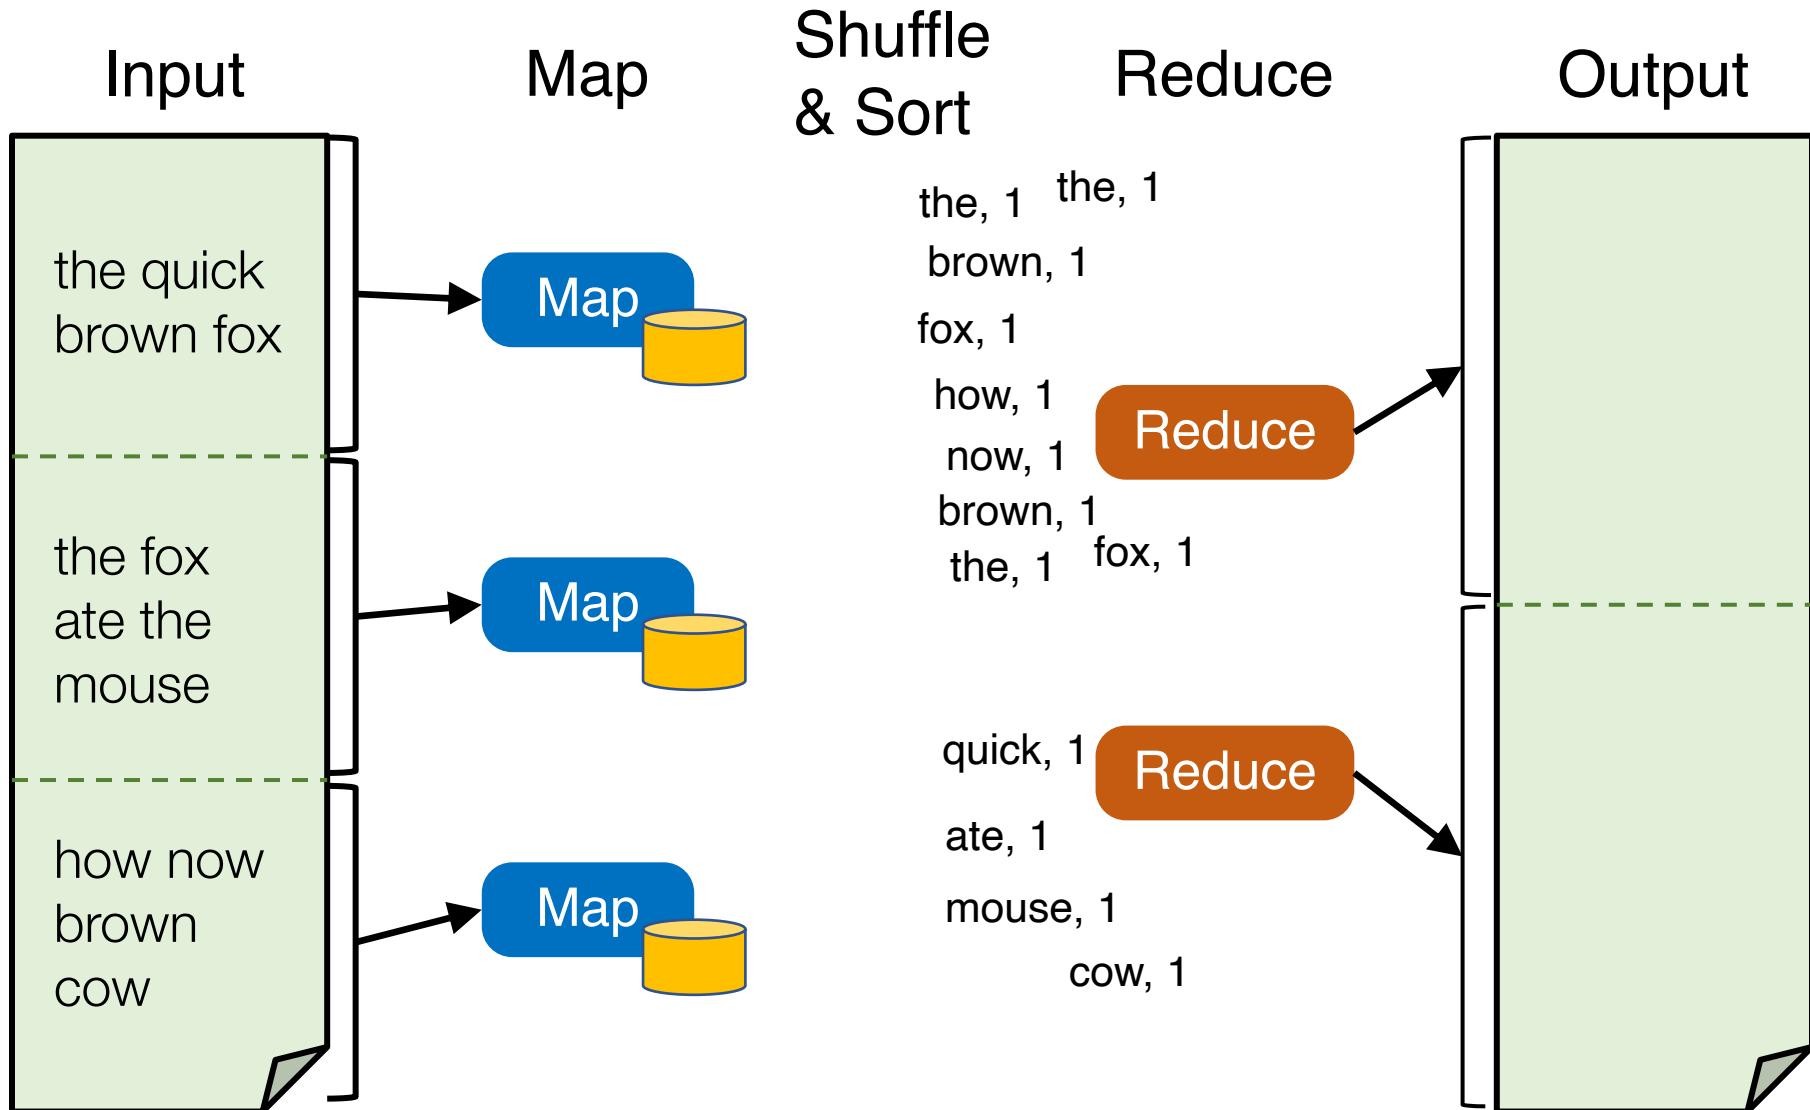
```
        results += ParseInt(v);
```

```
    Emit(AsString(result));
```

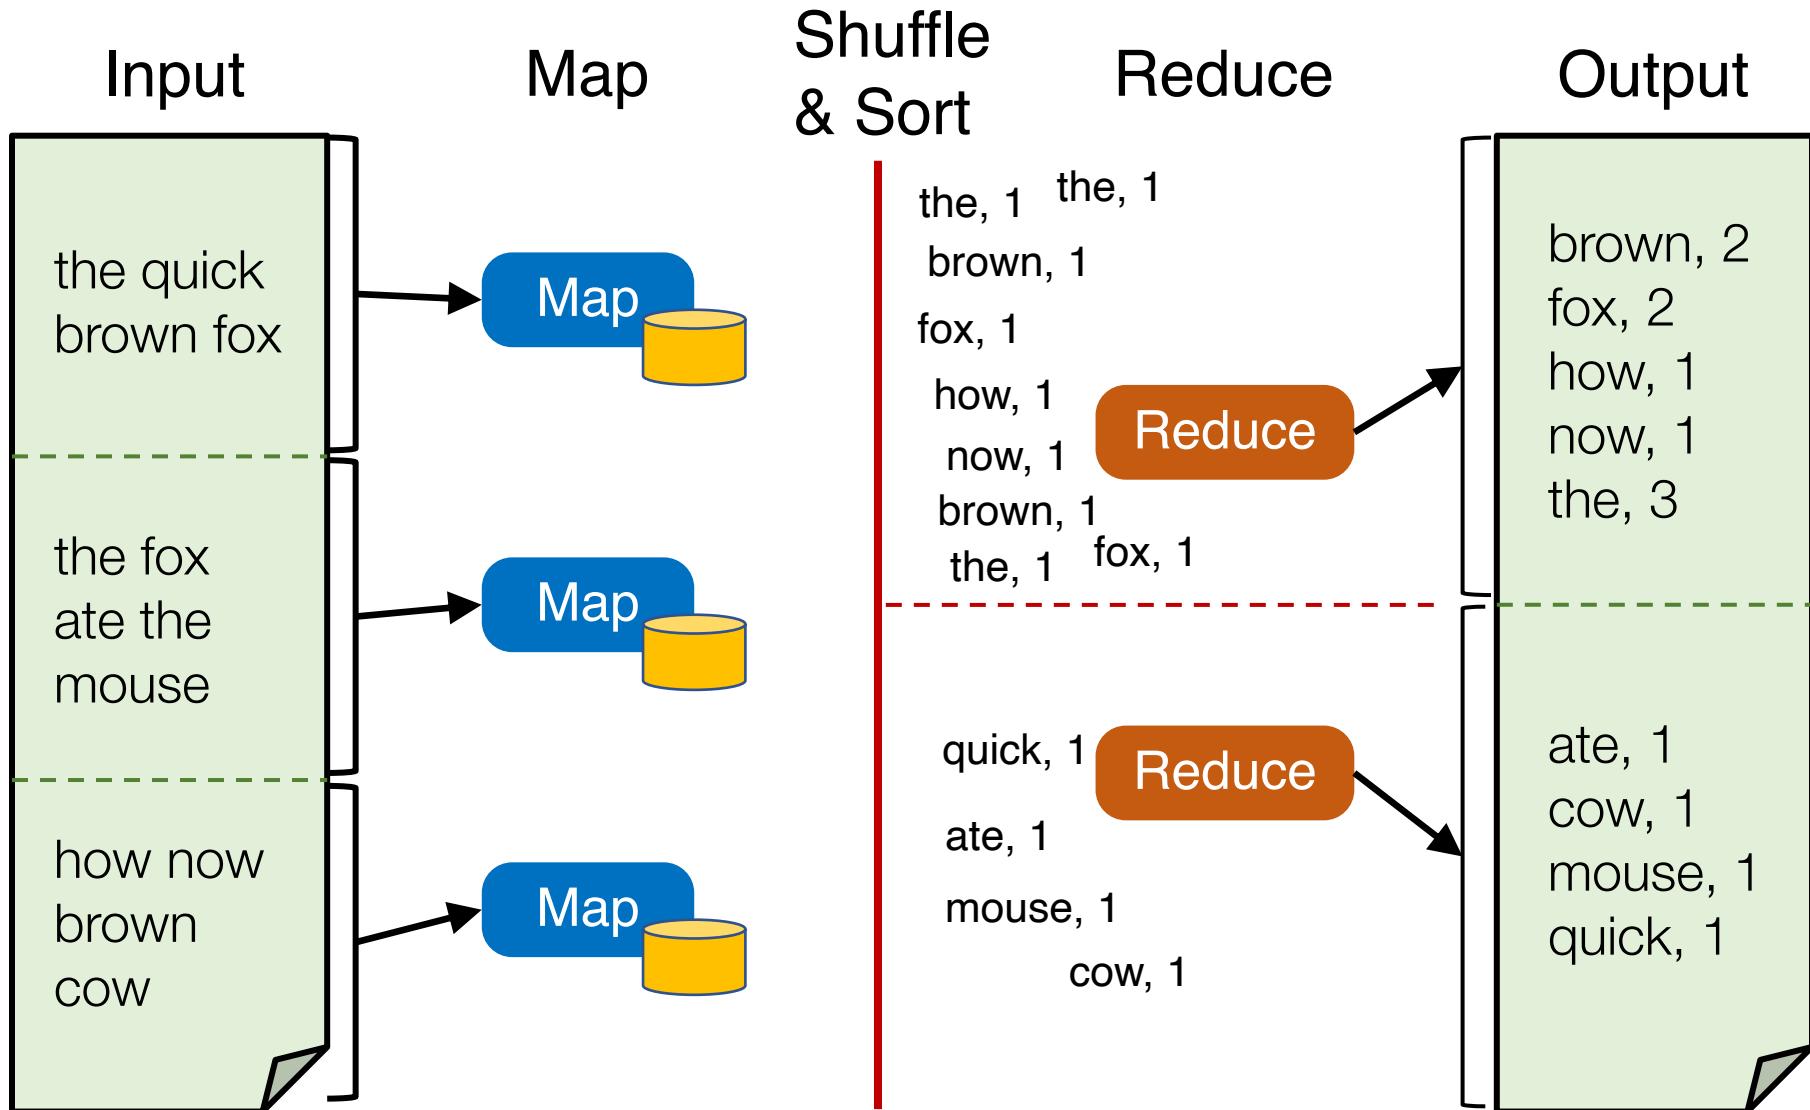
# Word Count execution



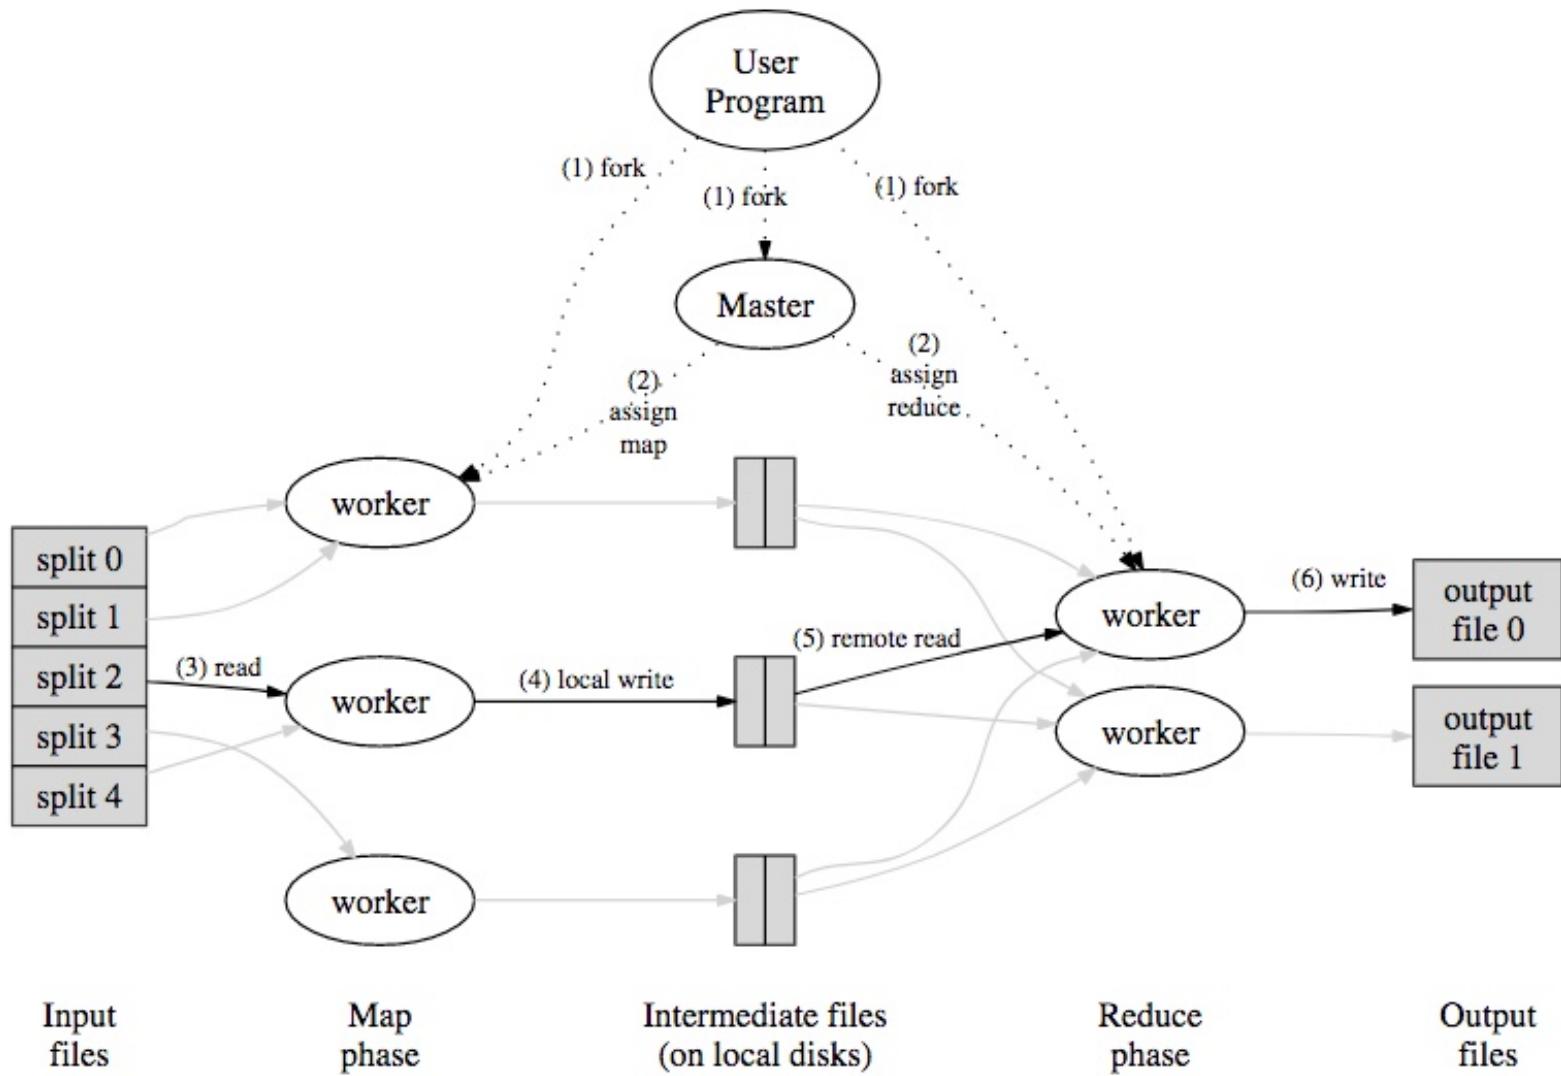
# Word Count execution



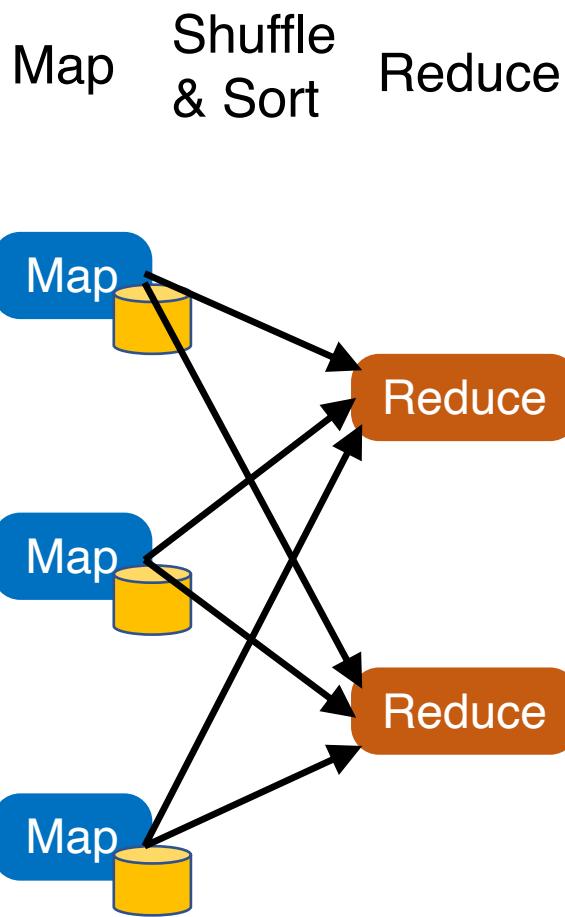
# Word Count execution



# MapReduce data flows



# MapReduce processes



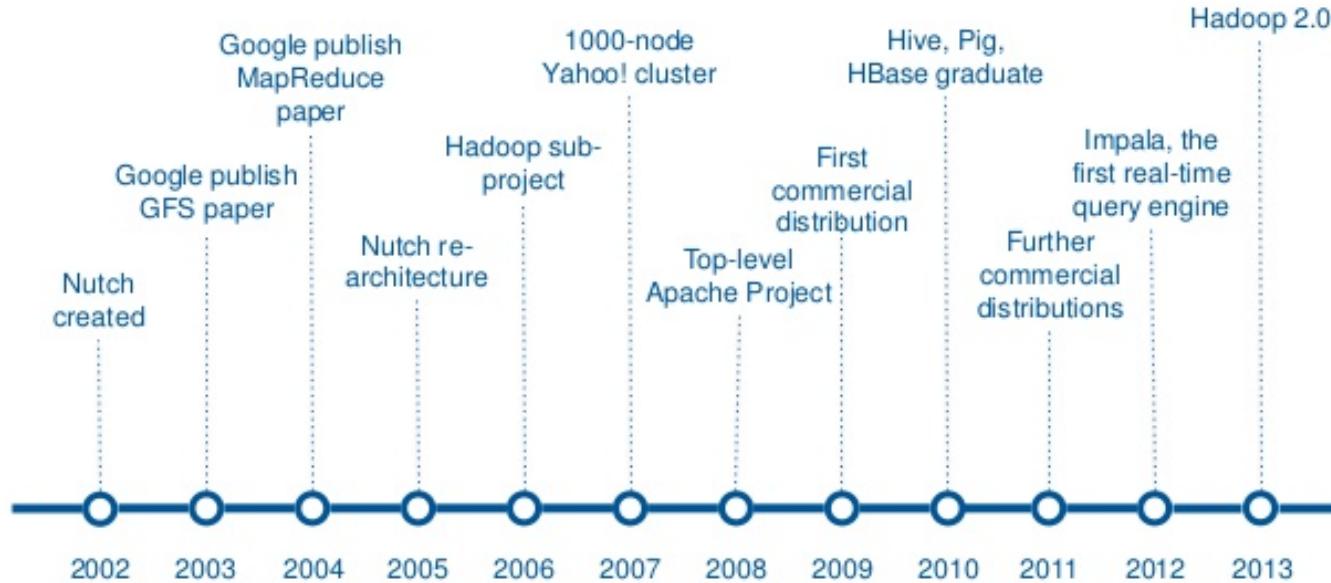
- Map workers write intermediate output to local disk, separated by partitioning. Once completed, tell master node
- Reduce worker told of location of map task outputs, pulls their partition's data from each mapper, execute function across data
- Note:
  - “All-to-all” shuffle b/w mappers and reducers
  - Written to disk (“materialized”) b/w each state

# Apache Hadoop

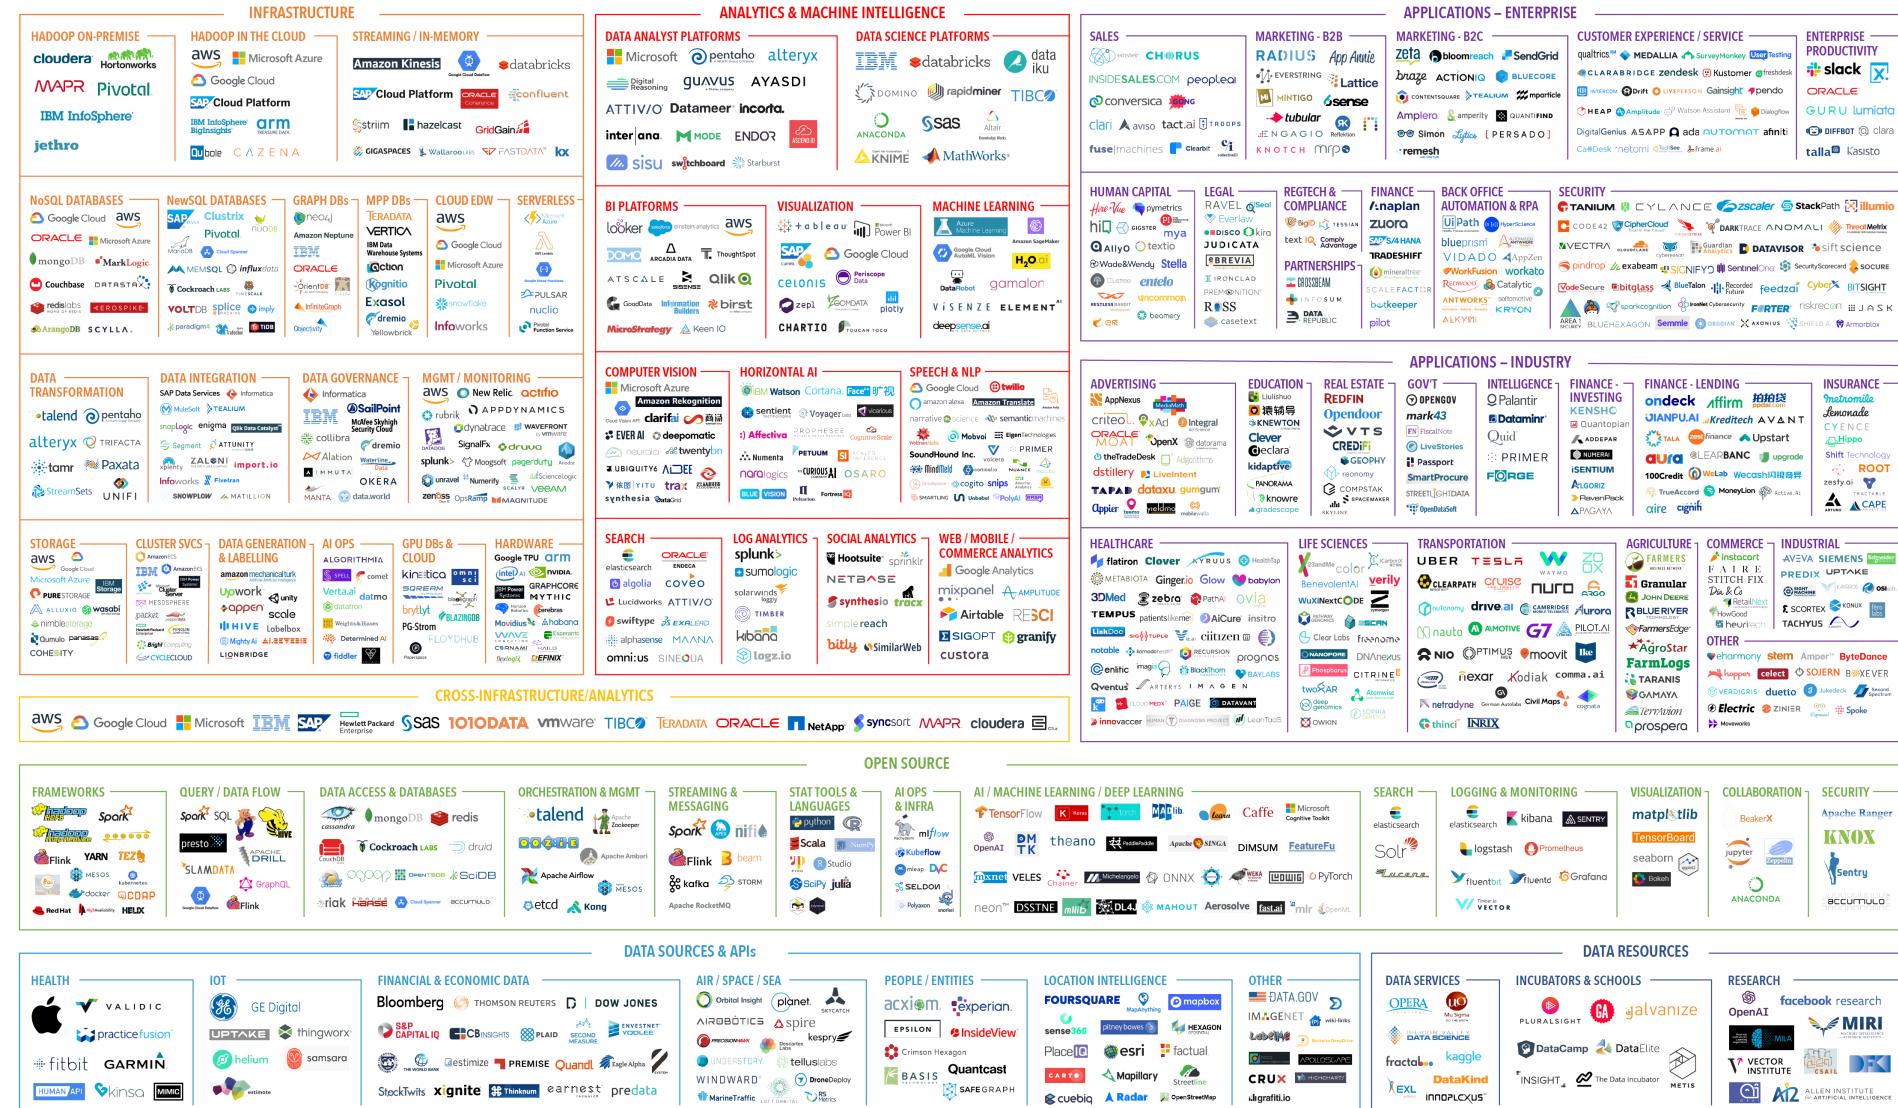


- An open-source implementation of Google's MapReduce framework
  - Hadoop MapReduce atop Hadoop Distributed File System (HDFS)

## A Brief History of Hadoop



DATA & AI LANDSCAPE 2019



July 16, 2019 - FINAL 2019 VERSION

© Matt Turck (@mattturck), Lisa Xu (@lisaxu92), & FirstMark (@firstmarkcap)

mattturck.com/data2019

FIRSTMARK  
EARLY STAGE VENTURE CAPITAL

# Go RPC

# Go RPCs

- Implementation in built-in library net/rpc

# Go RPCs

- Implementation in built-in library net/rpc
- Write stub receiver methods of the form
  - `func (t *T) MethodName(args T1, reply *T2) error`

# Go RPCs

- Implementation in built-in library net/rpc
- Write stub receiver methods of the form
  - `func (t *T) MethodName(args T1, reply *T2) error`
- Register receiver methods

# Go RPCs

- Implementation in built-in library net/rpc
- Write stub receiver methods of the form
  - `func (t *T) MethodName(args T1, reply *T2) error`
- Register receiver methods
- Create a listener (i.e., server) that accepts requests

# Writing a WordCount RPC server in Go

```
type WordCountServer struct {
    addr string
}

type WordCountRequest struct {
    Input string
}

type WordCountReply struct {
    Counts map[string]int
}
```

# Writing a WordCount RPC server in Go

```
type WordCountServer struct {
    addr string
}

type WordCountRequest struct {
    Input string
}

type WordCountReply struct {
    Counts map[string]int
}

func (*WordCountServer) Compute(
    request WordCountRequest,
    reply *WordCountReply) error {
    counts := make(map[string]int)
    input := request.Input
    tokens := strings.Fields(input)
    for _, t := range tokens {
        counts[t] += 1
    }
    reply.Counts = counts
    return nil
}
```

# Writing a WordCount RPC server in Go

```
type WordCountServer struct {
    addr string
}

type WordCountRequest struct {
    Input string
}

type WordCountReply struct {
    Counts map[string]int
}
```

```
func (*WordCountServer) Compute(
    request WordCountRequest,
    reply *WordCountReply) error {
    counts := make(map[string]int)
    input := request.Input
    tokens := strings.Fields(input)
    for _, t := range tokens {
        counts[t] += 1
    }
    reply.Counts = counts
    return nil
}
```

# Writing a WordCount RPC server in Go

```
func (server *WordCountServer) Listen() {
    rpc.Register(server)
    listener, err := net.Listen("tcp", server.addr)
    checkError(err)
    go func() {
        rpc.Accept(listener)
    }()
}
```

# Writing a WordCount RPC server in Go

```
func (server *WordCountServer) Listen() {
    rpc.Register(server)
    listener, err := net.Listen("tcp", server.addr)
    checkError(err)
    go func() {
        rpc.Accept(listener)
    }()
}
```

# Writing a WordCount RPC server in Go

```
func (server *WordCountServer) Listen() {
    rpc.Register(server)
    listener, err := net.Listen("tcp", server.addr)
    checkError(err)
    go func() {
        rpc.Accept(listener)
    }()
}
```

# Writing a WordCount RPC server in Go

```
func (server *WordCountServer) Listen() {
    rpc.Register(server)
    listener, err := net.Listen("tcp", server.addr)
    checkError(err)
    go func() {
        rpc.Accept(listener)
    }()
}
```

# WordCount client

```
func makeRequest(input string, serverAddr string) (map[string]int, error) {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    err = client.Call("WordCountServer.Compute", args, &reply)
    if err != nil {
        return nil, err
    }
    return reply.Counts, nil
}
```

# WordCount client

```
func makeRequest(input string, serverAddr string) (map[string]int, error) {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    err = client.Call("WordCountServer.Compute", args, &reply)
    if err != nil {
        return nil, err
    }
    return reply.Counts, nil
}
```

# WordCount client

```
func makeRequest(input string, serverAddr string) (map[string]int, error) {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    err = client.Call("WordCountServer.Compute", args, &reply)
    if err != nil {
        return nil, err
    }
    return reply.Counts, nil
}
```

# WordCount client

```
func makeRequest(input string, serverAddr string) (map[string]int, error) {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    err = client.Call("WordCountServer.Compute", args, &reply)
    if err != nil {
        return nil, err
    }
    return reply.Counts, nil
}
```

# WordCount client-server

```
func main() {
    serverAddr := "localhost:8888"
    server := WordCountServer{serverAddr}
    server.Listen()
    input1 := "hello I am good hello bye bye bye bye good night hello"
    wordcount, err := makeRequest(input1, serverAddr)
    checkError(err)
    fmt.Printf("Result: %v\n", wordcount)
}
```

# WordCount client-server

```
func main() {
    serverAddr := "localhost:8888"
    server := WordCountServer{serverAddr}
    server.Listen()
    input1 := "hello I am good hello bye bye bye good night hello"
    wordcount, err := makeRequest(input1, serverAddr)
    checkError(err)
    fmt.Printf("Result: %v\n", wordcount)
}
```

```
Result: map[hello:3 I:1 am:1 good:2 bye:4 night:1]
```

# Is this synchronous or asynchronous?

```
func makeRequest(input string, serverAddr string) (map[string]int, error)
{
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    err = client.Call("WordCountServer.Compute", args, &reply)
    if err != nil {
        return nil, err
    }
    return reply.Counts, nil
}
```

# Making client asynchronous

```
func makeRequest(input string, serverAddr string) chan Result {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}

    return ch
}
```

# Making client asynchronous

```
func makeRequest(input string, serverAddr string) chan Result {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    ch := make(chan Result)
    go func() {
        err := client.Call("WordCountServer.Compute", args, &reply)
        if err != nil {
            ch <- Result{nil, err} // something went wrong
        } else {
            ch <- Result{reply.Counts, nil} // success
        }
    }()
    return ch
}
```

# Making client asynchronous

```
func makeRequest(input string, serverAddr string) *Call {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    return client.Go("WordCountServer.Compute", args, &reply, nil)
}
```

# Making client asynchronous

```
func makeRequest(input string, serverAddr string) *Call {
    client, err := rpc.Dial("tcp", serverAddr)
    checkError(err)
    args := WordCountRequest{input}
    reply := WordCountReply{make(map[string]int)}
    return client.Go("WordCountServer.Compute", args, &reply, nil)
}
```

```
call := makeRequest(...)
<-call.Done
checkError(call.Error)
handleReply(call.Reply)
```